BIOS 721|TOPIC 1: OBJECTS IN R – PART 2

Outline

- Last time we covered …
  - Mode of objects available in R
    - Numeric / Character / Logical
  - Class of objects available in R
    - Vector / Matrix / Array / Data Frame / List
  - Functions used to create objects in R

- This time we will cover …
  - Basic manipulations performed on objects in R
    - Indexing
    - Applying operators
    - Applying simple functions

Indexing

- In Topic 1 Part 1, we discussed how to create several types/modes of objects in R.

- Once an object has been created, either by the user or as the output from an R function, the user will often need to select specific elements of that object for further manipulation.
  - This is known as ‘indexing’.

Indexing – Motivating Example

- For example, consider the data frame below. Suppose we wanted to create another data frame that subsets to all male subjects who were exposed to treatment A.
  - How could we get R to identify these subjects using code? ⟶ Indexing!

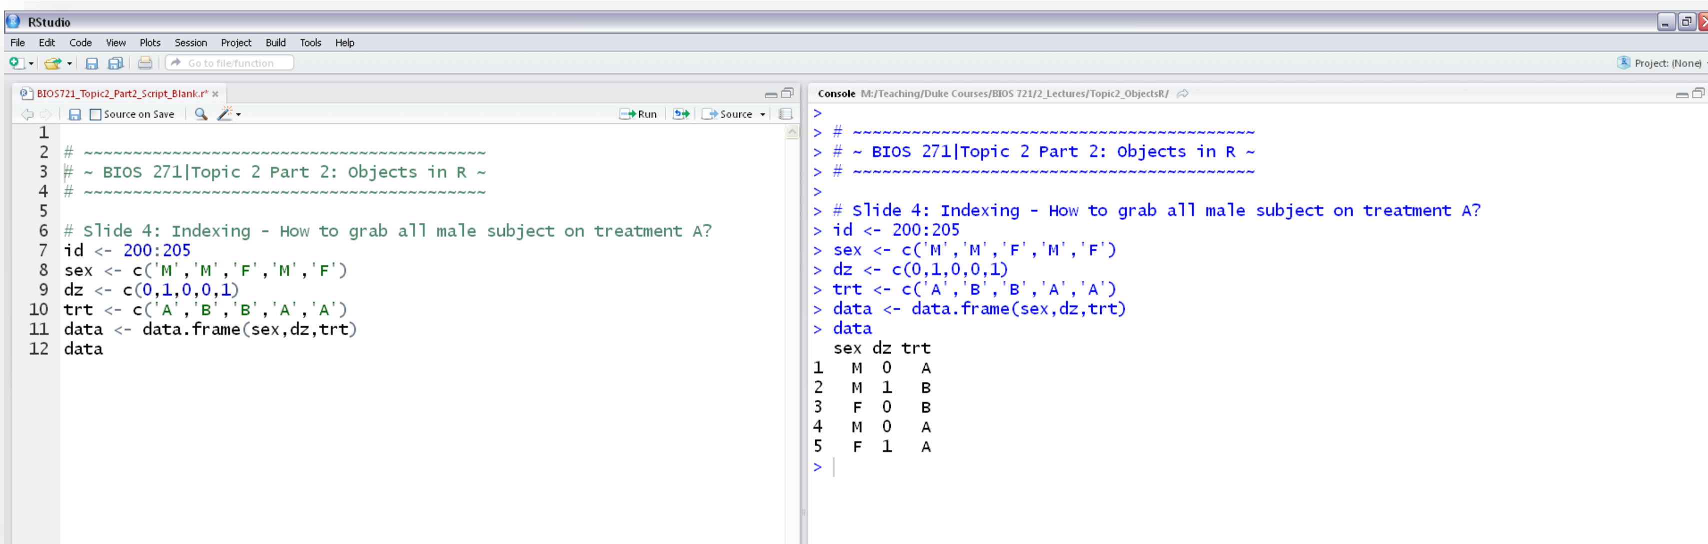

Indexing

- In general, indexing allows users to modify objects by selecting or removing specified elements from the original object to created a new object.

- In R, users can index elements in object using their numerical indices (e.g. row and column NUMBER in a matrix) or by their dimension names if available (e.g. row and column NAME in a matrix).

- Let's see examples of how indexing works for the following objects. Although the syntax is slightly different for each object, there is a pattern to the syntax.
  - Vectors
  - Matrices
  - Data Frames
  - Will discuss indexing other objects (Arrays and Lists) later in the course

Indexing

- To select elements in a vector, indicate the position of the elements you want to retain or drop within square brackets next the object name.
  - Indexing Vectors: `object[element(s) position]`
  - **Example|** Consider the vector `v <- c(1,5,4,0,9)`

```
v[2]
v[c(1,4)]
v[1:3]
v[7]
v[-3]
```

Indexing

- Indexing Vectors

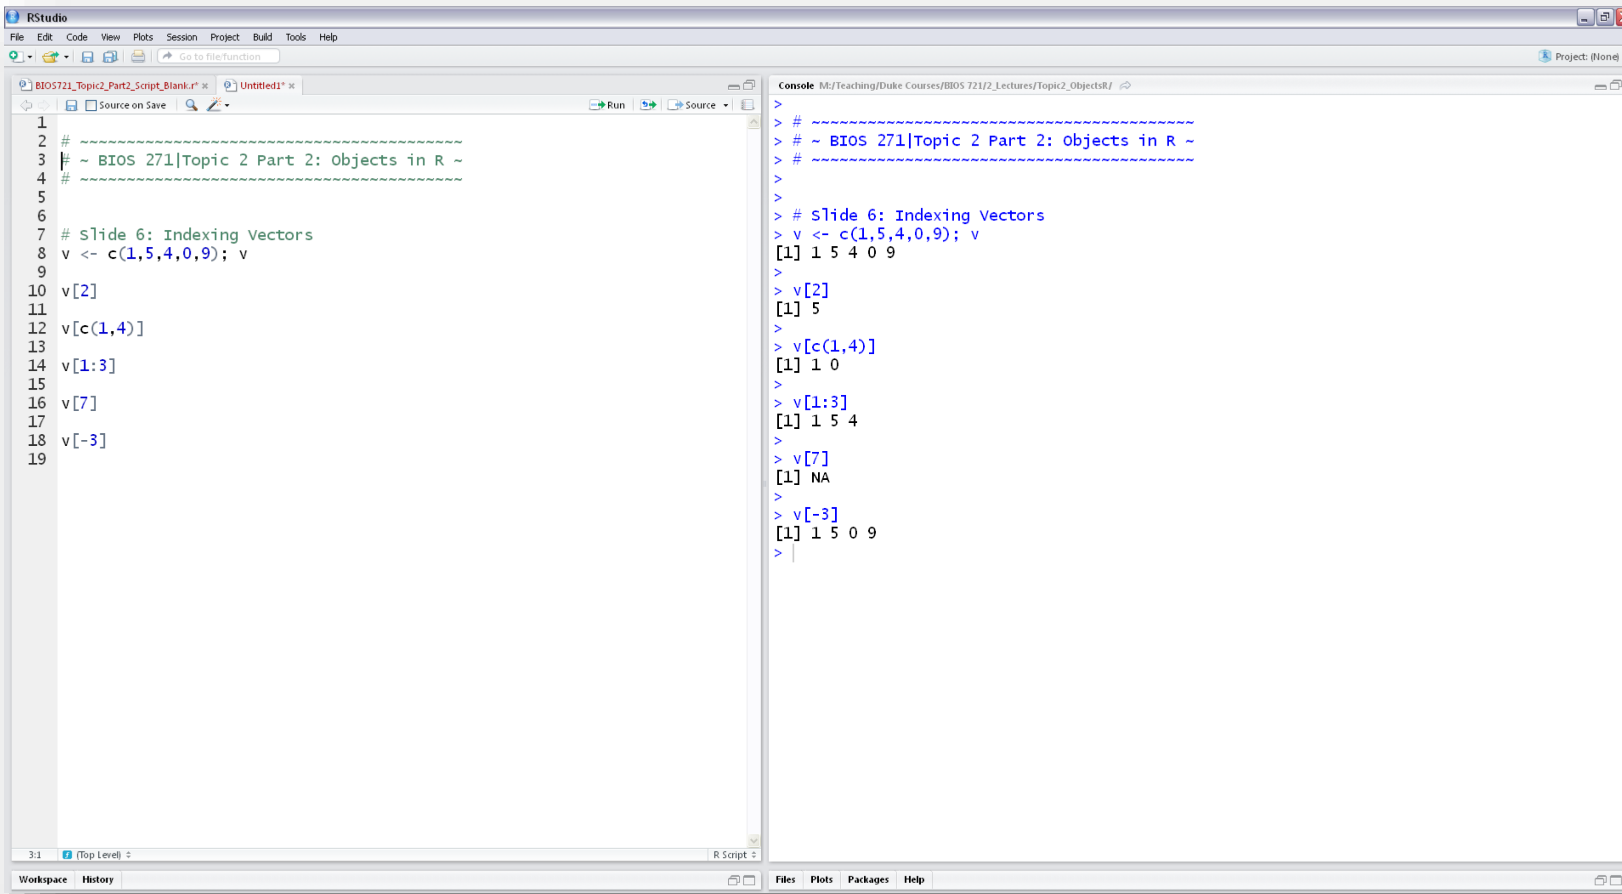

Indexing

- To select elements in a matrix, indicate the position of the elements you want to retain or drop within square brackets next the object name.
  - Indexing Matrices: `object[row(s),column(s)]`
  - **Example|** Consider the matrix `M <- matrix(1:6,3,2)`

```
M[2,1]
M[1,2]
M[,1]
M[2:3,]
M[,-1]
```

Indexing

- To select elements in a matrix, indicate the position of the elements you want to retain or drop within square brackets next the object name.
  - Indexing Matrices: `object[row(s),column(s)]`
  - **Example|** Consider the matrix `M <- matrix(1:6,3,2)`

```
M[2,1]
M[1,2]
M[,1]
M[2:3,]
M[,-1]
```

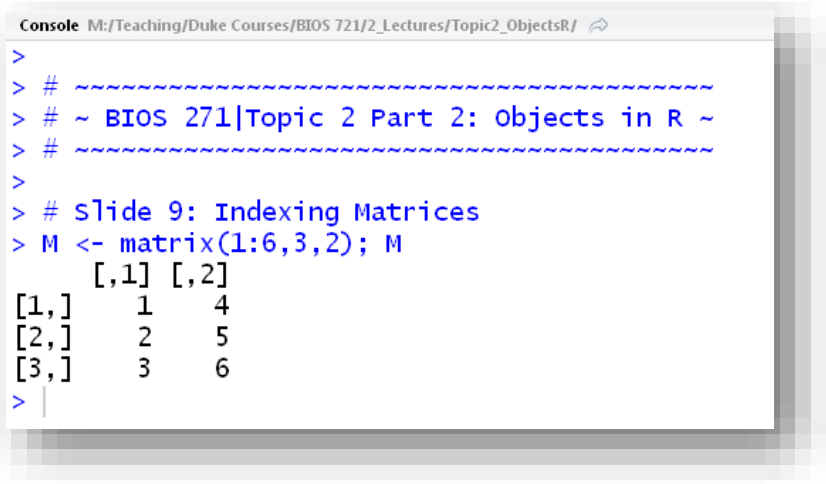

Indexing

- Indexing Matrices

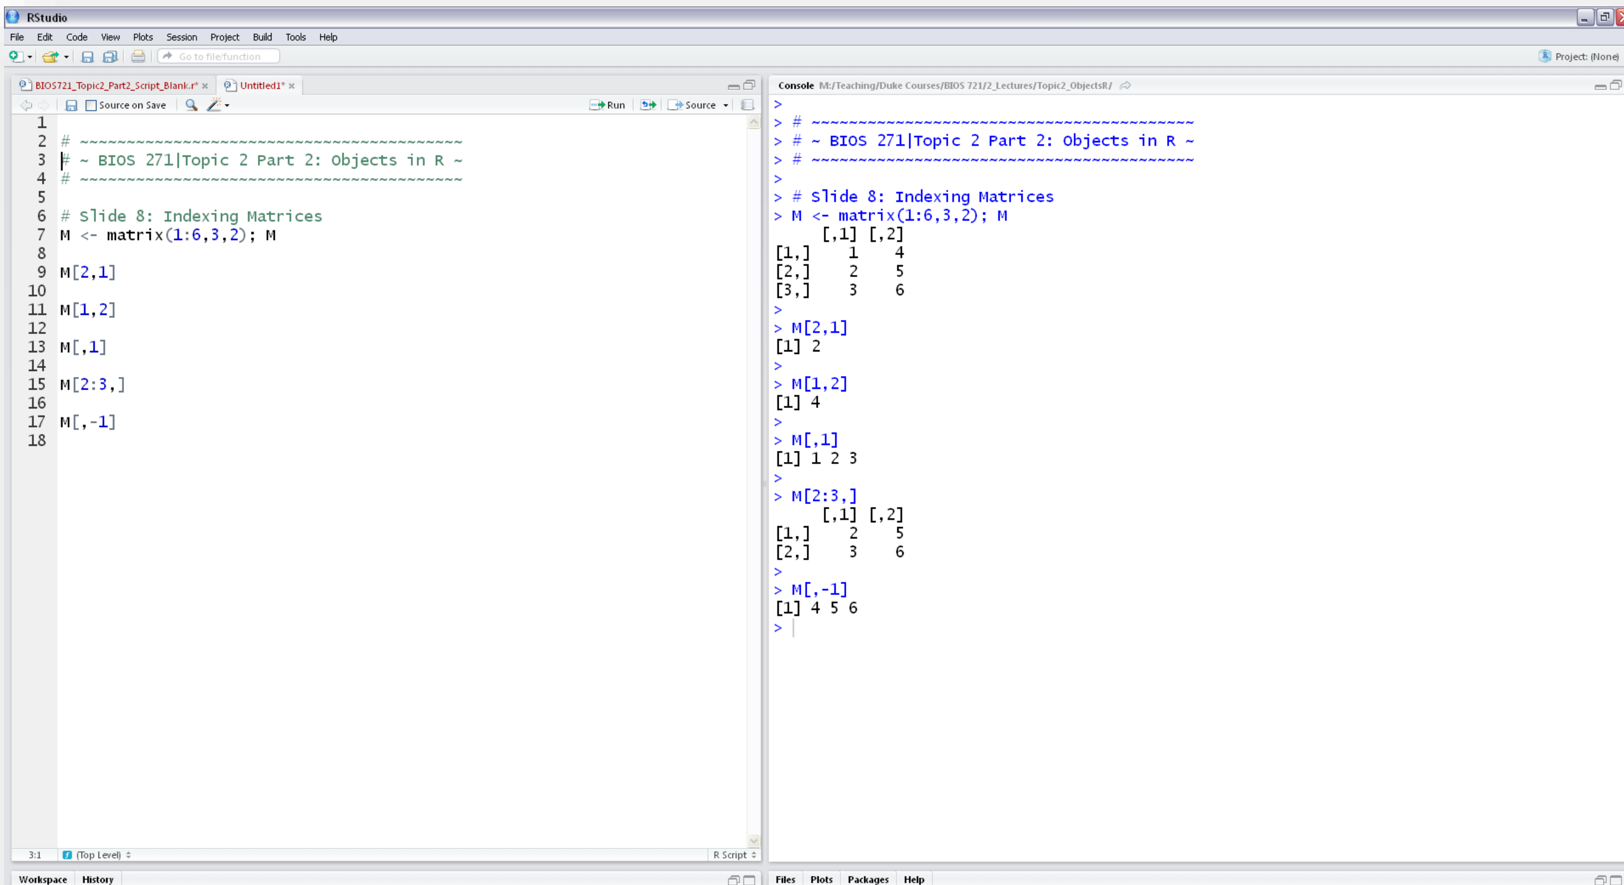

Indexing

- Can also index matrices using row and column names to identify the elements that you would like to retain.
  - Indexing Matrices: `object["rownames","columnnames"]`
  - **Example|** Consider the matrix

```
M <- matrix(1:4,2,2,byrow=TRUE)
rownames(M) <- c('r1','r2')
colnames(M) <- c('C1','C2')

M['r2','C2']
M['r2','c2']
M['r1',]
M[,'C1']
```

Indexing

- Can also index matrices using row and column names to identify the elements that you would like to retain.
  - Indexing Matrices: `object["rownames","columnnames"]`

  - **Example|** Consider the matrix

```
M <- matrix(1:4,2,2,byrow=TRUE)
rownames(M) <- c('r1','r2')
colnames(M) <- c('C1','C2')

M['r2','C2']
M['r2','c2']
M['r1',]
M[,'C1']
```

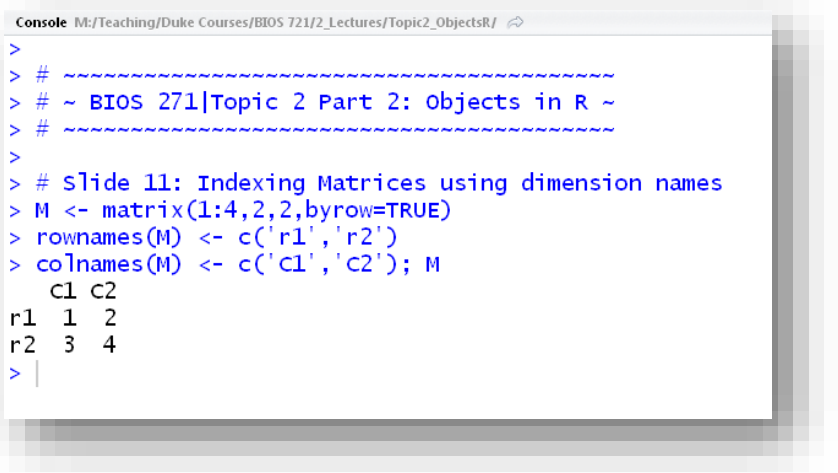

Indexing

- Indexing Matrices

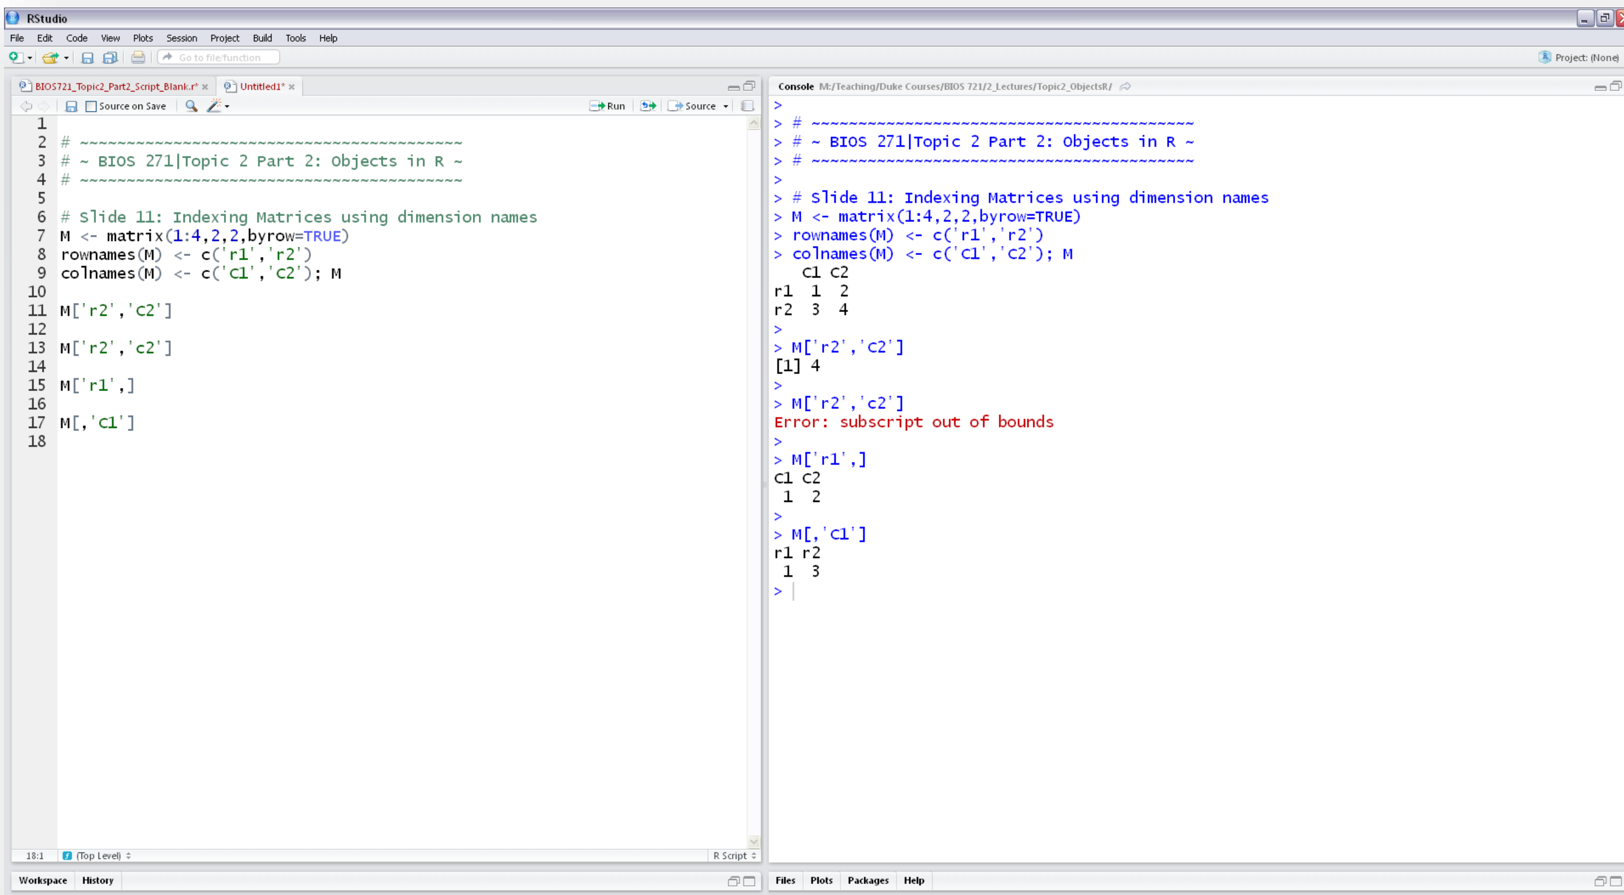

Indexing

- To select elements in a data frame, indicate the position of the elements you want to retain or drop within square brackets next the object name.
  - Indexing Data Frames: `object[row(s),column(s)]`

  - Example| Consider the data frame

```
id <- 200:205
sex <- c('M','M','F','M','F')
dz <- c(0,1,0,0,1)
trt <- c('A','B','B','A','A')
D <- data.frame(sex,dz,trt)

D[1,3]
D[2:4,2]
D[,1]
```

Indexing

- To select elements in a data frame, indicate the position of the elements you want to retain or drop within square brackets next the object name.
  - Indexing Data Frame: `object[row(s),column(s)]`

  - **Example|** Consider the data frame

```
id <- 200:205
sex <- c('M','M','F',
dz <- c(0,1,0,0,1)
trt <- c('A','B','B',
D <- data.frame(sex,d

D[1,3]
D[2:4,2]
D[,1]
```

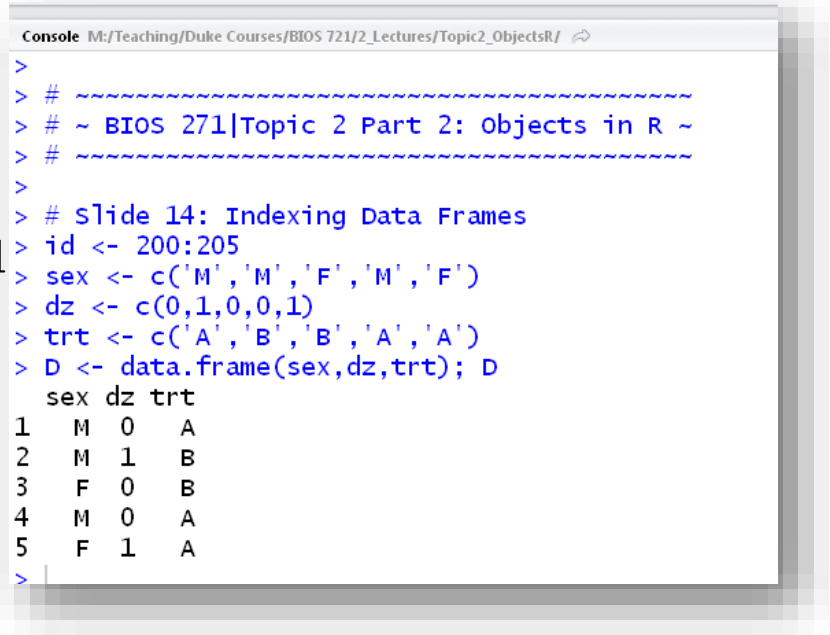

Indexing

- Indexing Data Frames

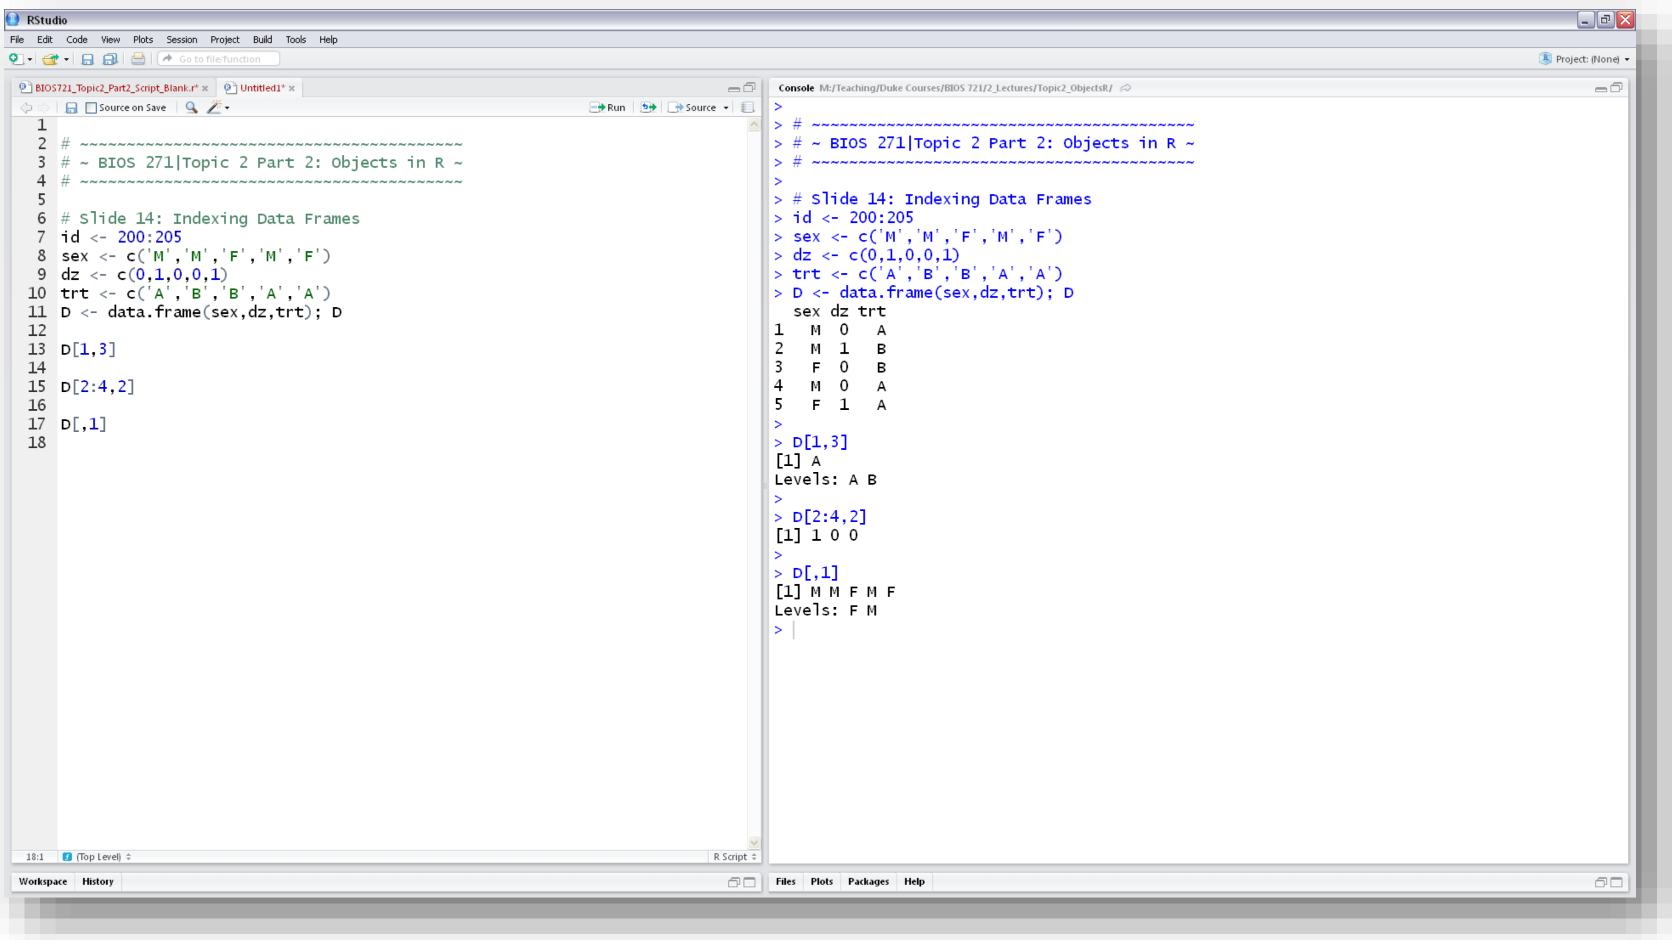

Indexing

- Can also index data frames using the column names to identify the columns that you would like to retain.
  - Indexing Data Frames: `object[row(s),"columnnames"]`
    `object$columnname[row(s)]`
  - **Example|** Consider the data frame

```
id <- 200:205
sex <- c('M','M','F','M','F')
dz <- c(0,1,0,0,1)
trt <- c('A','B','B','A','A')
D <- data.frame(sex,dz,trt)

D[1,'sex']
D[,'trt']
D$trt
D$dz[c(1,4)]
```

Indexing

- Can also index data frames using the column names to identify the columns that you would like to retain.
  - Indexing Data Frames: `object[row(s),"columnnames"]`
    `object$columnname[row(s)]`
  - **Example|** Consider the data frame

```
id <- 200:205
sex <- c('M','M','F',
dz <- c(0,1,0,0,1)
trt <- c('A','B','B',
D <- data.frame(sex,d

D[1,'sex']
D[,'trt']
D$trt
D$dz[c(1,4)]
```

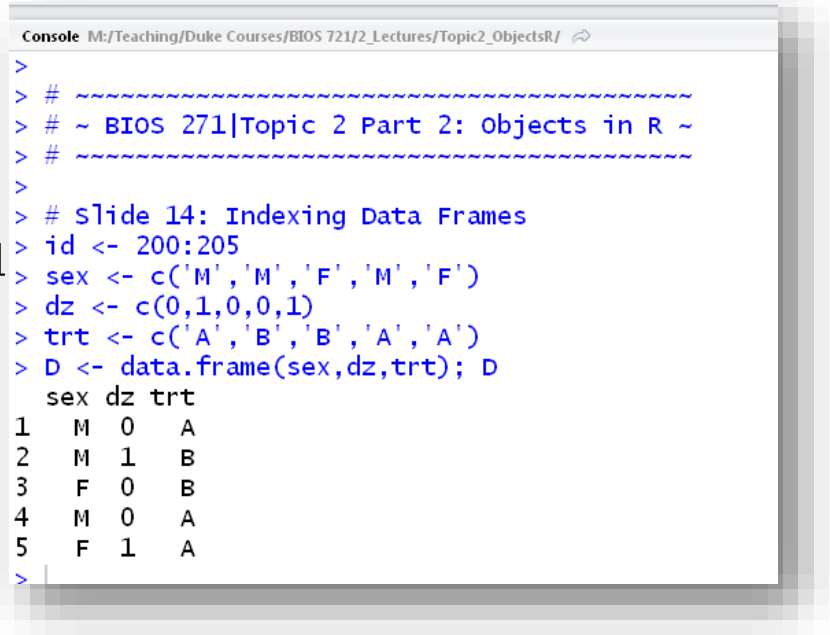

Indexing

- Indexing Data Frames

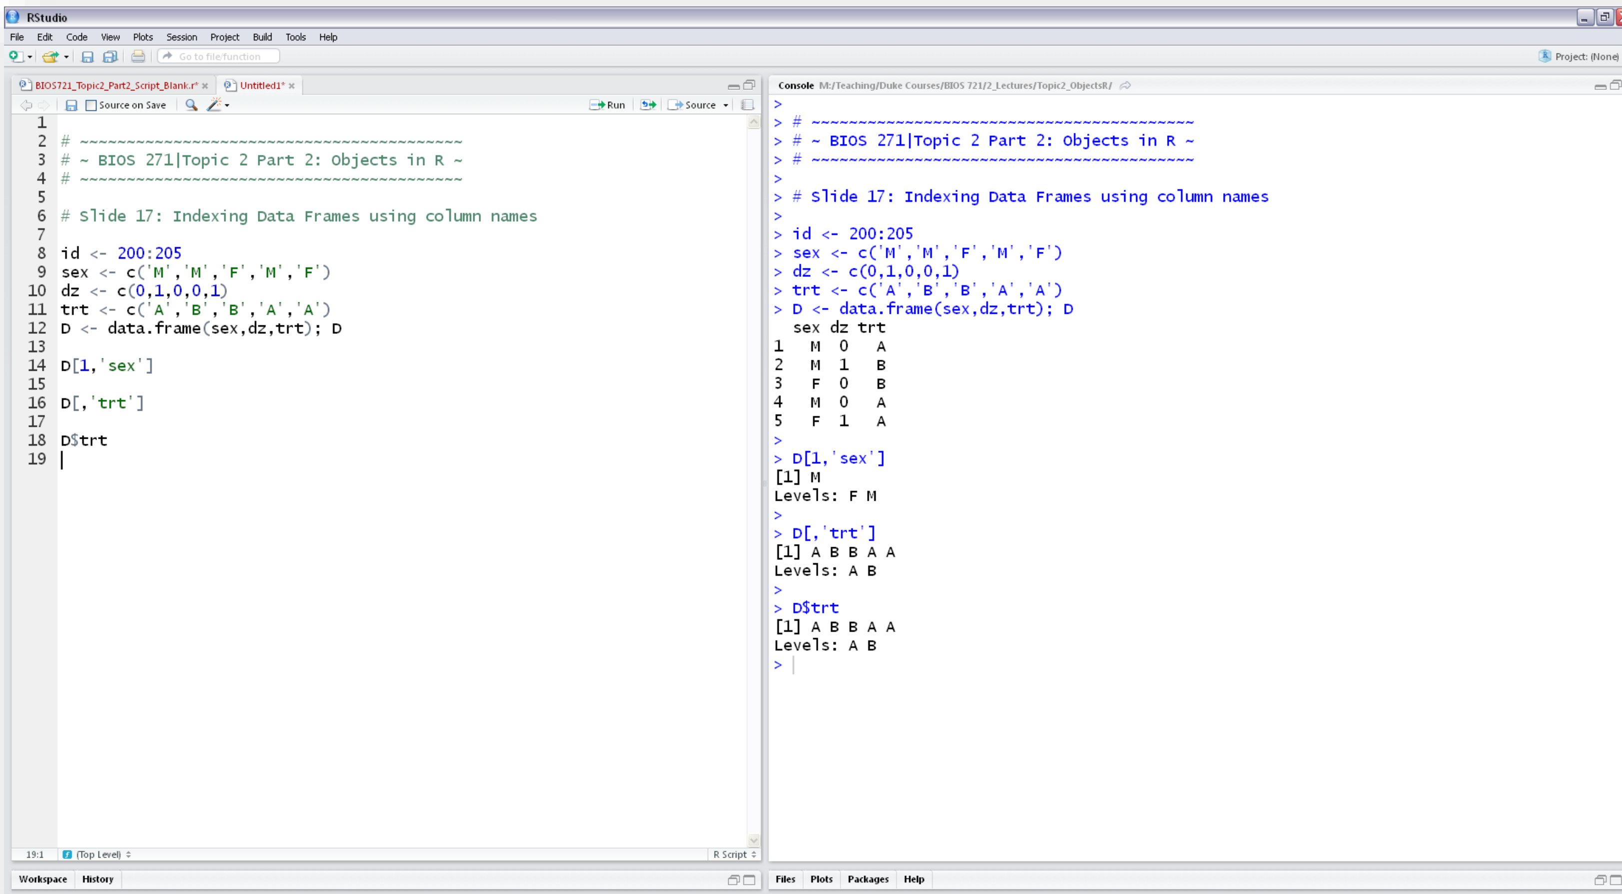

Indexing – Motivating Example

- Recall this example?
  - Consider the data frame below. Suppose we wanted to create another data frame that subsets to all male subjects who were exposed to treatment A.
  - How could we get R to identify these subjects using code?
    - Indexing!
    - Using indexing, could create a data frame called sub that contains this subset of the data.

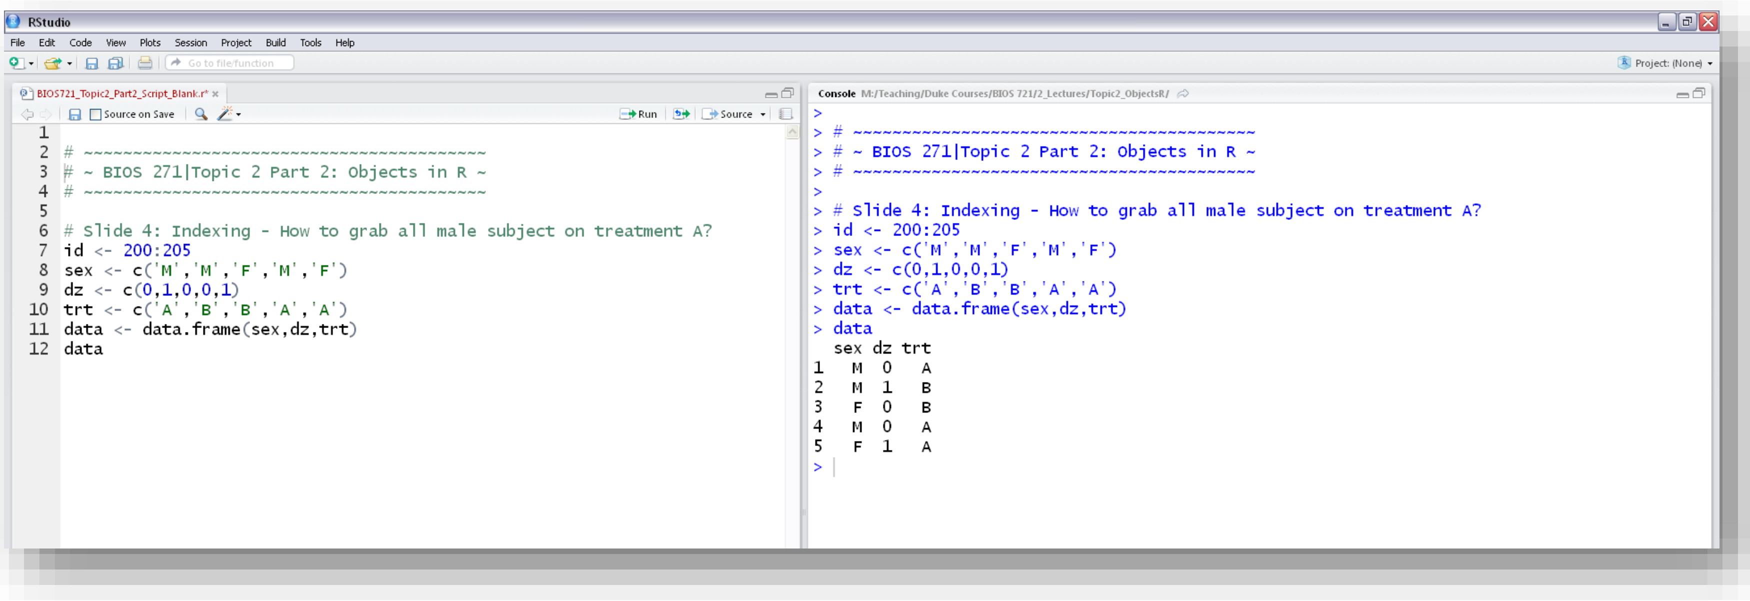

Indexing – Motivating Example

- Recall this example?
  - For solution – come to class!

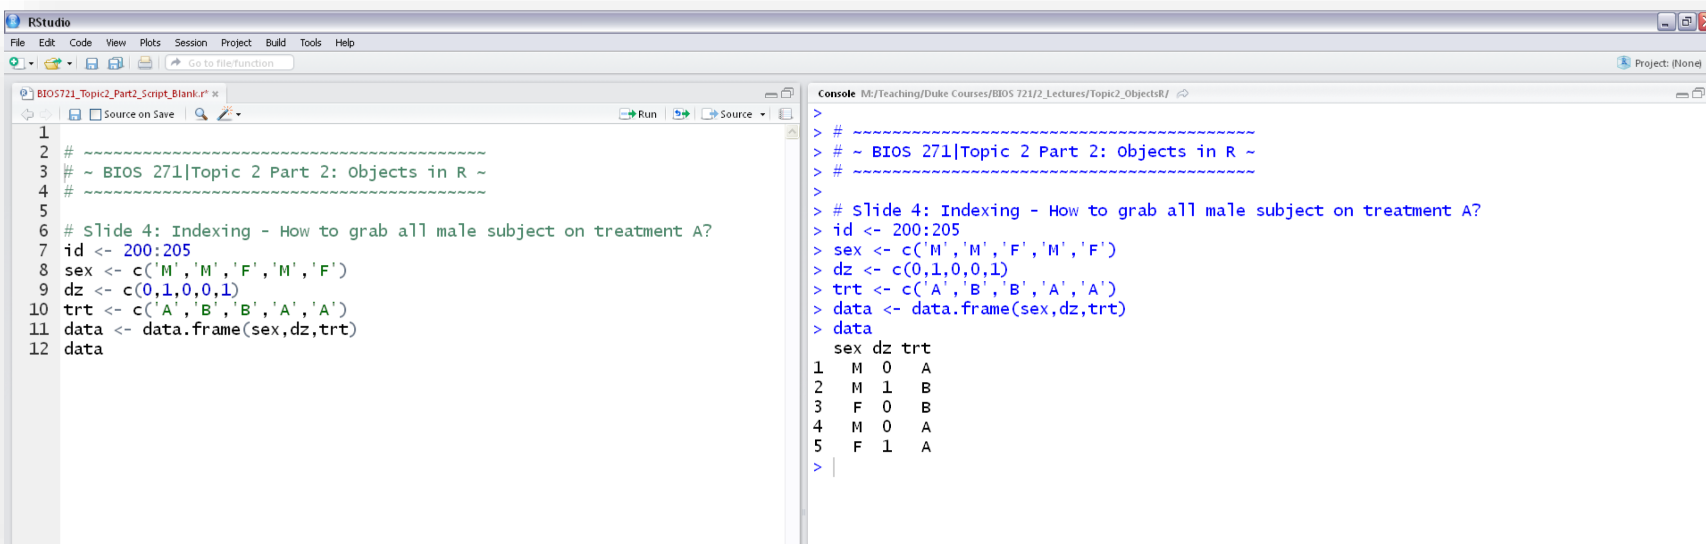

Indexing – Motivating Example

- Recall this example?
  - To create the data frame sub on the previous slide, we had to “hard code” indexing – ok for this example because the data set is small, but what if there were thousands on observations?
  - We would like to use code to <u>automate</u> identifying these subjects
    - Need to use indexing in conjunction with R operators and functions!

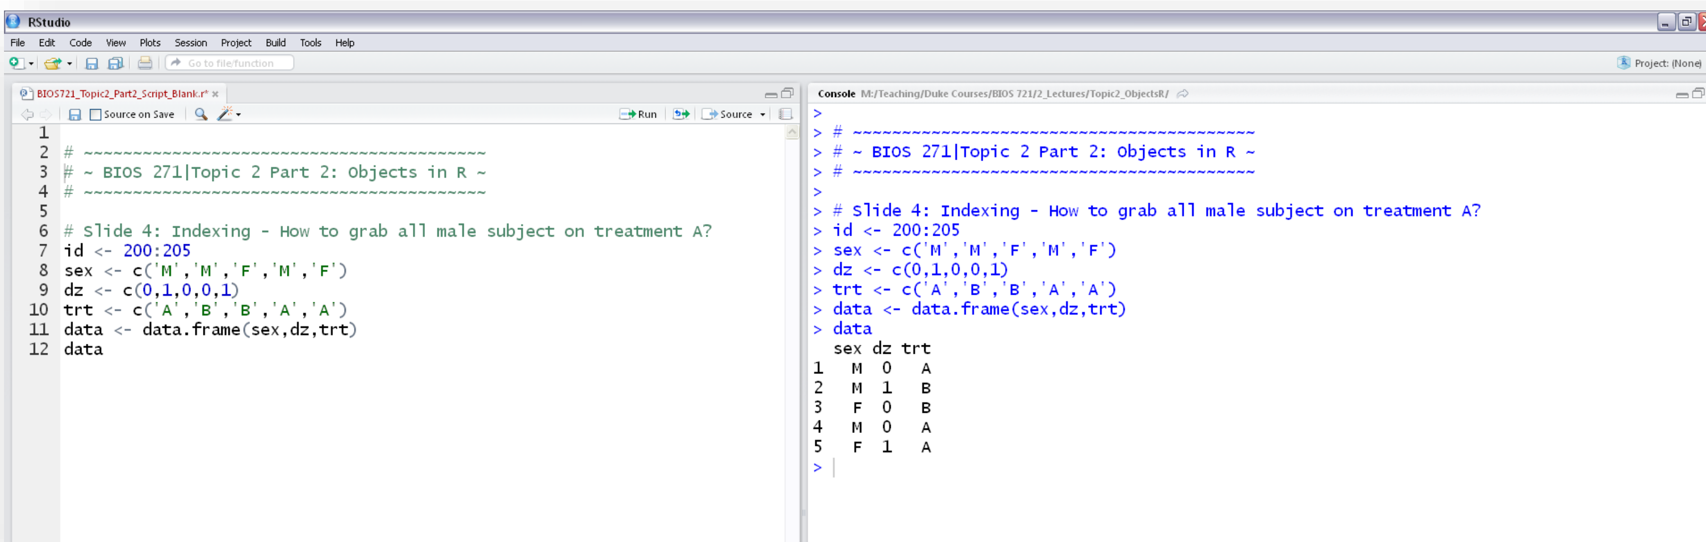

R Operators vs. R Functions

- R Operators
  - Set of supported language constructs (i.e. symbols) that perform operations on programming objects.
  - For example, + for addition

- R Functions
  - Packaged set of commands designed to perform a specific task.
  - Can be simple like `mean(x, …)` to calculate the mean of a data object x.
  - Can also be complex like `lm(y~x, …)` to estimate the linear relationship between the variables y and x.
    - For now, we will focus on simple R functions that everyone should have in their toolbox … actually not sure how you would program without knowing these functions.

R Operators

- + → addition
- − → subtraction
- * → multiplication
- / → division
- ^ → raise to a power
- %% → remainder
- %*% → matrix mult

- < → less than?
- > → greater than?
- <= → less than or equal?
- >= → greater than or equal to?
- ! → not
- == → equal to?
- != → not equal to?
- & → and
- | → or

R Operators

- + ⟶ addition
- − ⟶ subtraction
- * ⟶ multiplication
- / ⟶ division
- ^ ⟶ raise to a power
- %% ⟶ remainder
- %*% ⟶ matrix mult

Mathematical Operators (return #s)
vs.
Logical Operators (return TRUE/FALSE)

- < ⟶ less than?
- > ⟶ greater than?
- <= ⟶ less than or equal?
- >= ⟶ greater than or equal to?
- ! ⟶ not
- == ⟶ equal to?
- != ⟶ not equal to?
- & ⟶ and
- | ⟶ or

R Operators

- + → addition
- − → subtraction
- * → multiplication
- / → division
- ^ → raise to a power
- %% → remainder
- %*% → matrix mult

Mathematical Operators (return #s)
vs.
Logical Operators (return TRUE/FALSE)

- < → less than?
- > → greater than?
- <= → less than or equal?
- >= → greater than or equal to?
- ! → not
- == → equal to?
- != → not equal to?
- & → and
- | → or

R Operators

- These operators can be applied to any* object!
  - Will execute them element-wise (except for `%*%`)
    - That is, if applied to vector, will return a vector.
    - I.e. R is a 'vectorized' language – really useful trick
  - * Any compatible object (i.e. can't add two character values)

  - **Example|** Consider the vector `v <- 1:5`

```
v+2
v*v
v%*%v
v>3
v!=2
```

R Operators

- R applies operators element-wise

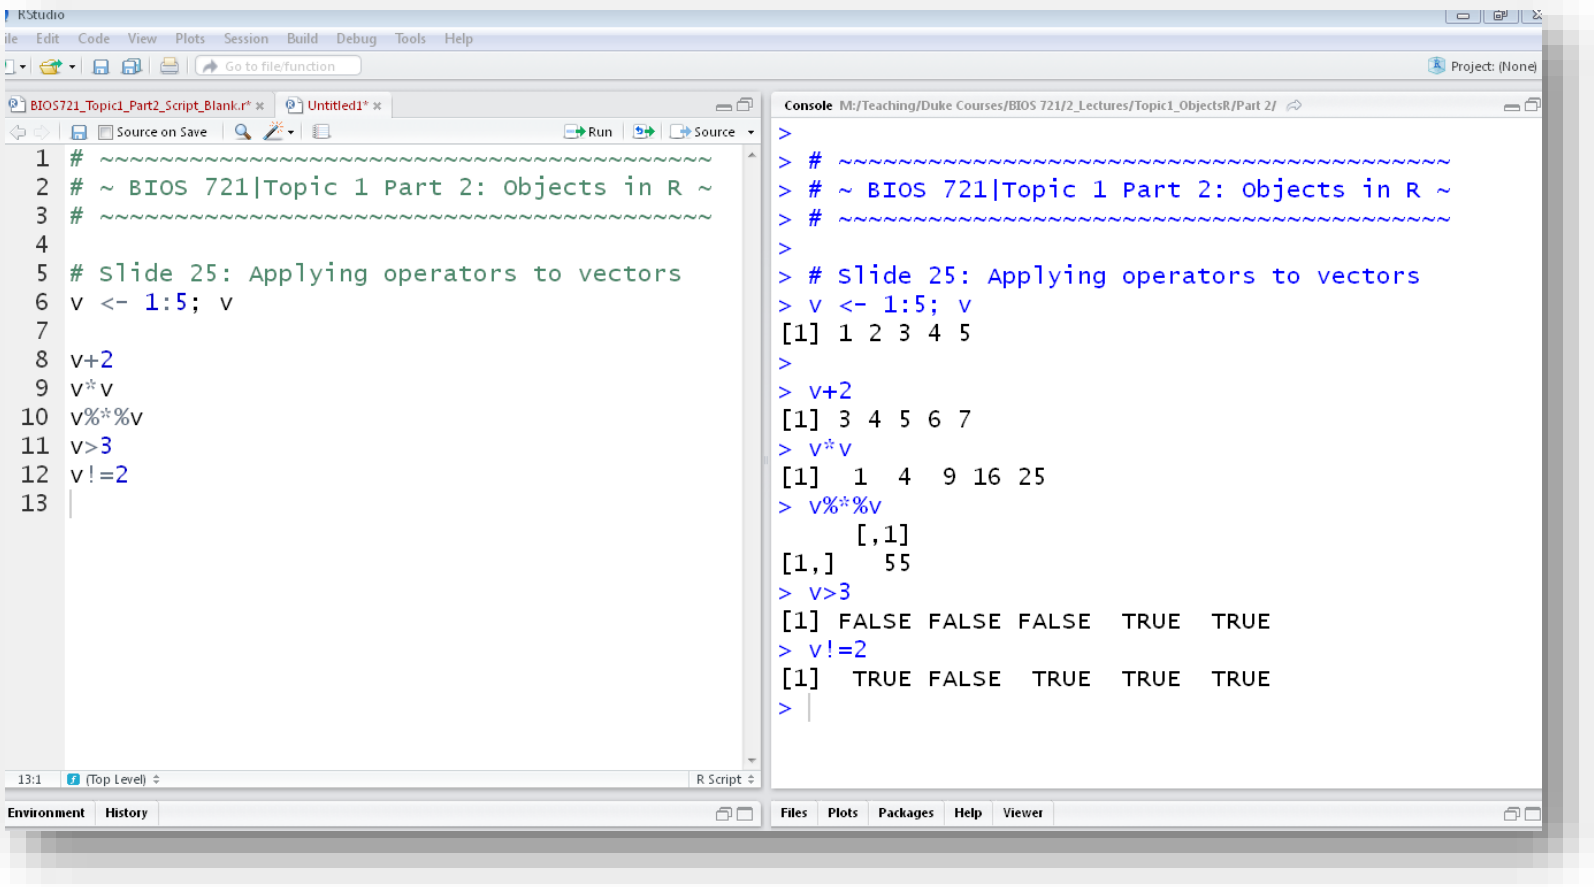

R Operators

- R applies operators element-wise

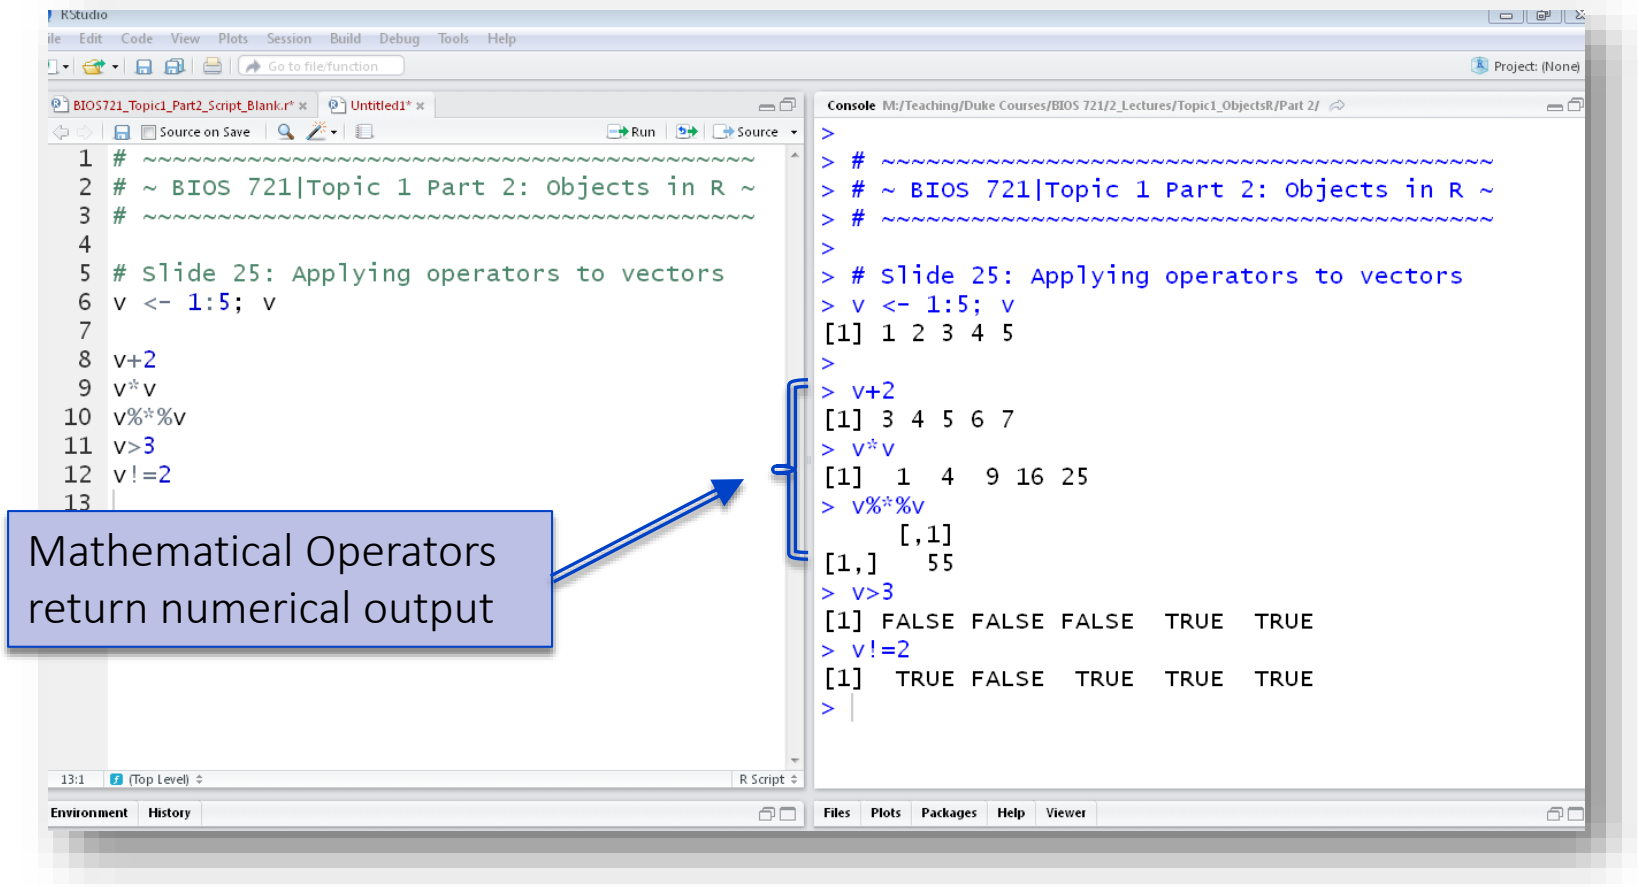

R Operators

- R applies operators element-wise

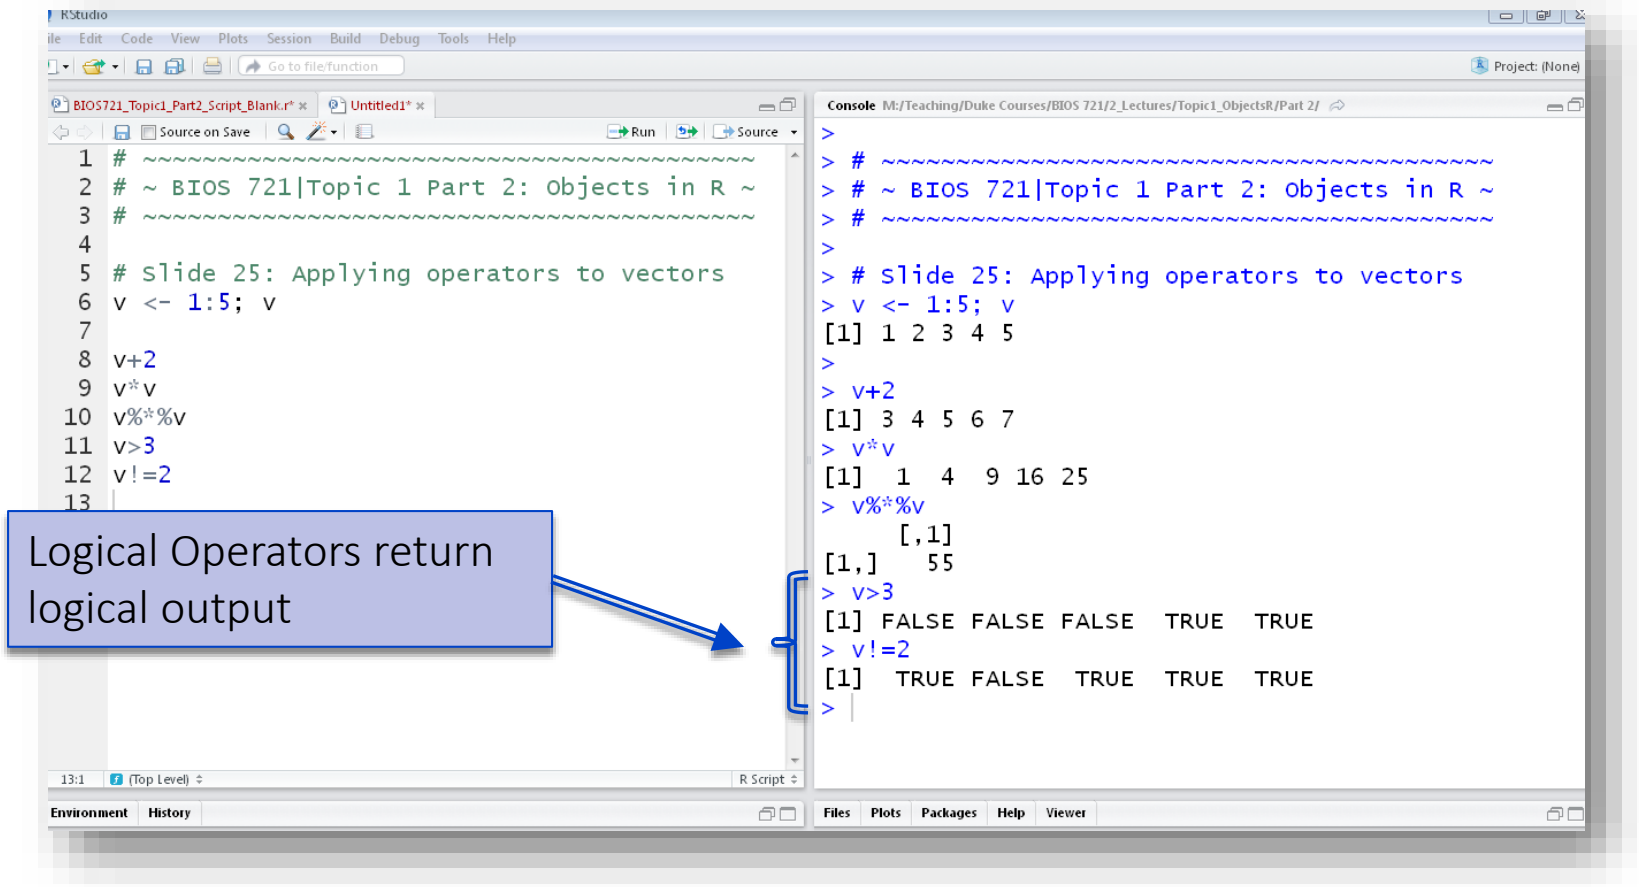

R Operators

- R applies operators element-wise

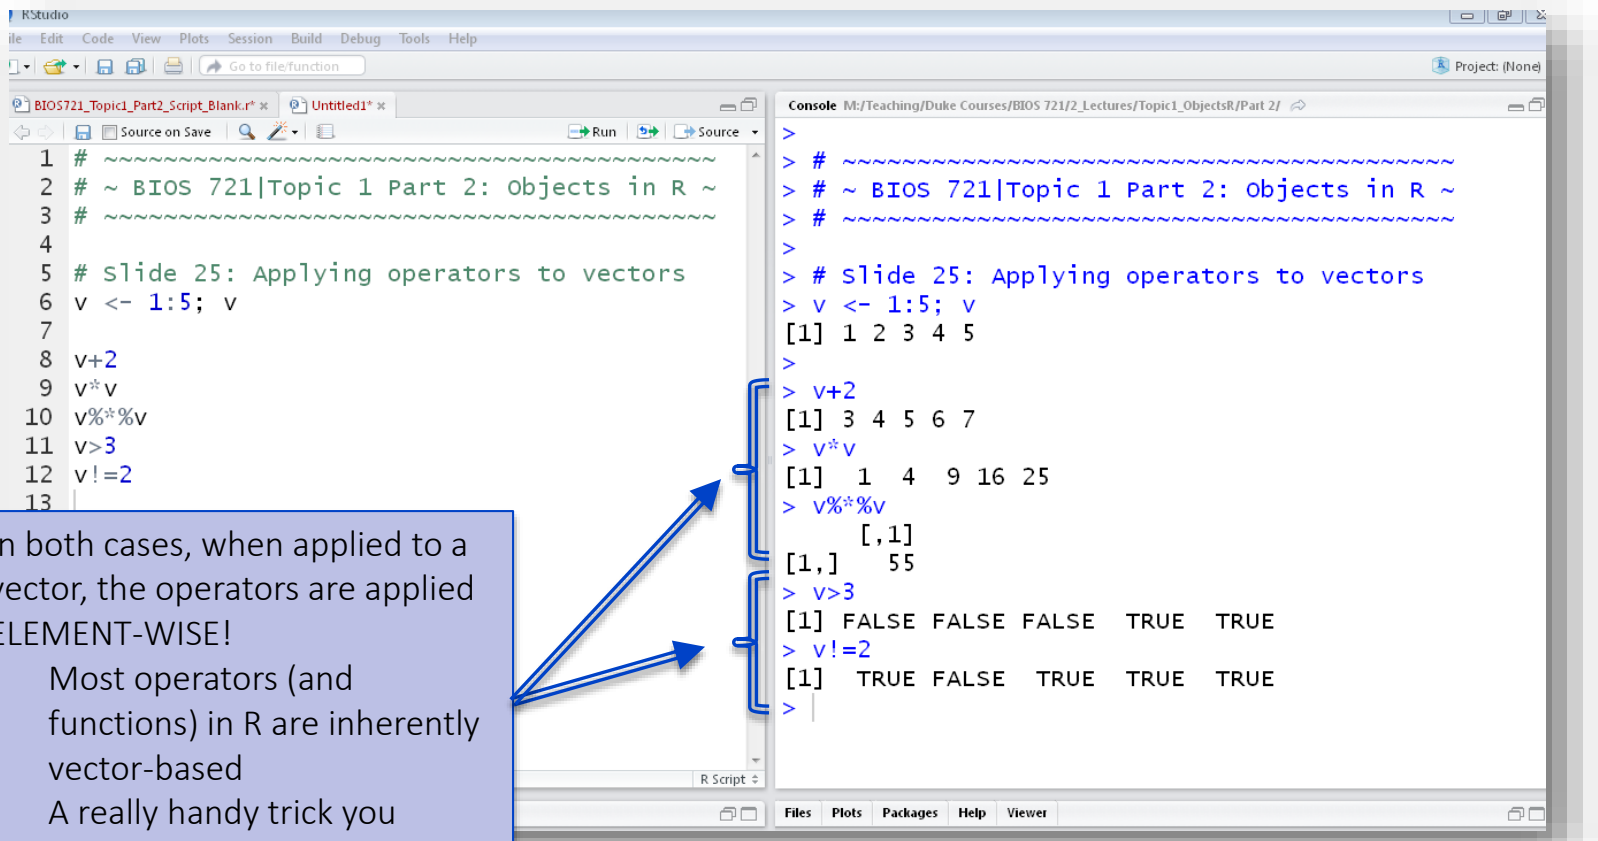

In both cases, when applied to a vector, the operators are applied ELEMENT-WISE!

- Most operators (and functions) in R are inherently vector-based
- A really handy trick you should take advantage of

R Operators

- The results are the same for matrices.
  - That is, will execute them element-wise (except for %*%)
    - I.e. R is a 'vectorized' language – really useful trick
    - Makes sense – a matrix is just an organized vector …
  - **Example|** Consider the matrices

```
M1 <- matrix(1:4,2,2)
M2 <- diag(rep(1,2))

M1/2
M1-M2
M1==3
```

R Operators

- The results are the same for matrices.
  - That is, will execute them element-wise (except for %*%)
    - I.e. R is a 'vectorized' language – really useful trick
    - Makes sense – a matrix is just an organized vector …
  - **Example|** Consider the matrices

```
M1 <- matrix(1:4,
M2 <- diag(rep(1,

M1/2
M1-M2
M1==3
```

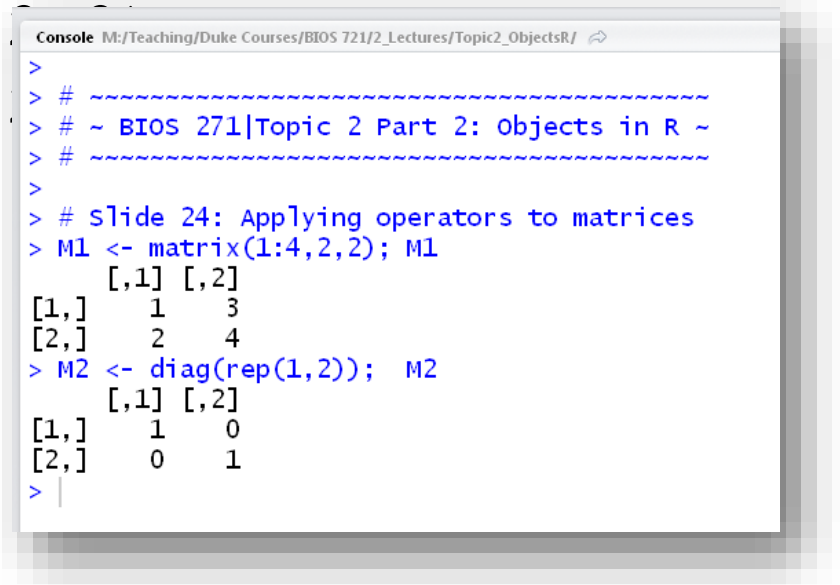

R Operators

- R applies operators element-wise

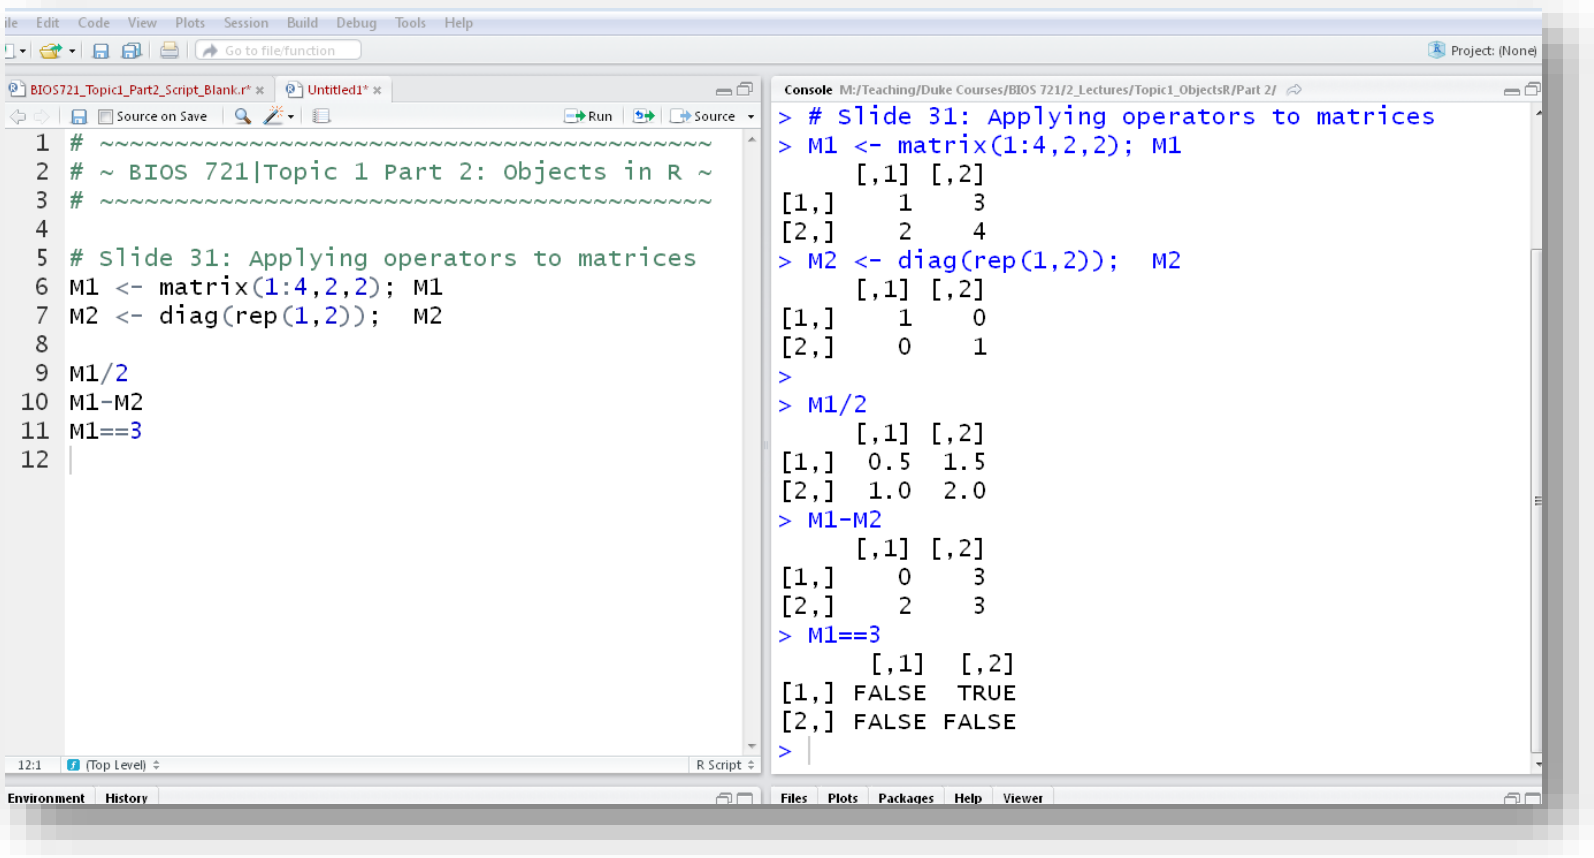

Indexing – Motivating Example

- Recall this example?
  - Consider the data frame below. Suppose we wanted to create another data frame that subsets to all male subjects who were exposed to treatment A.
  - How can operators be used to automate the indexing needed to create a data frame called sub that contains this subset of the data?

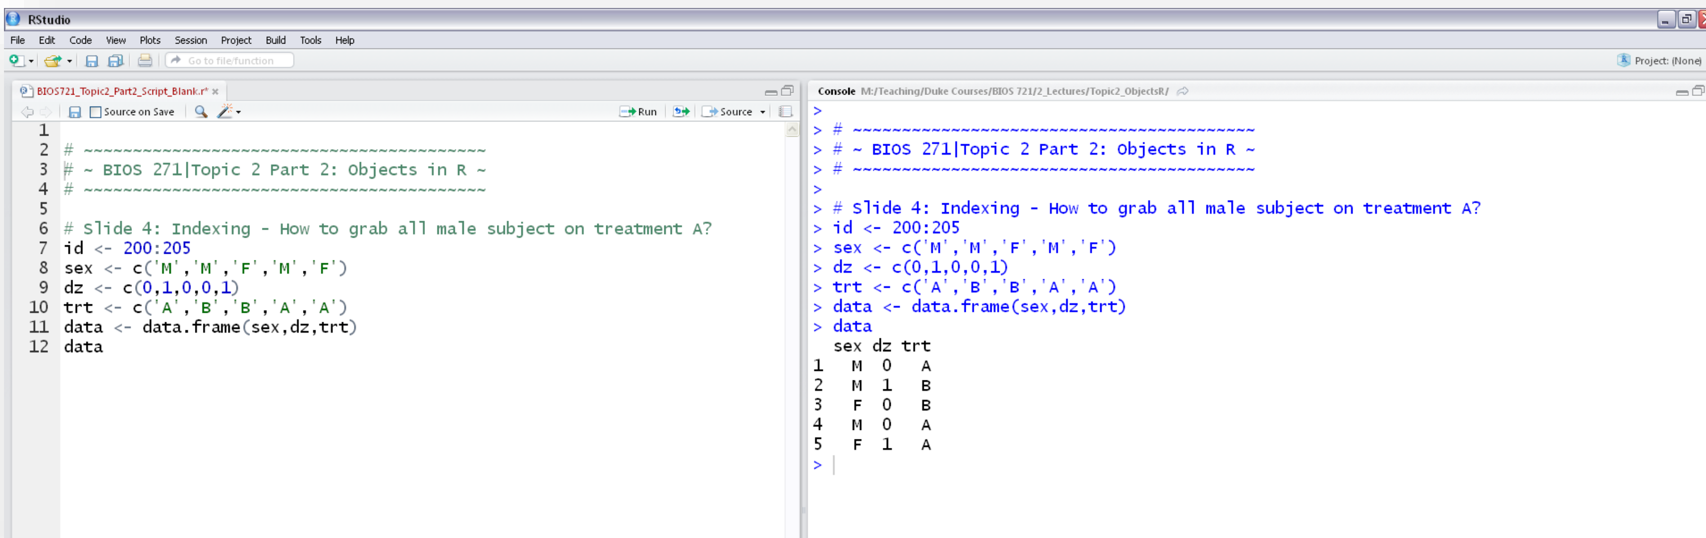

Indexing – Motivating Example

- Recall this example?
  - For solution – come to class!

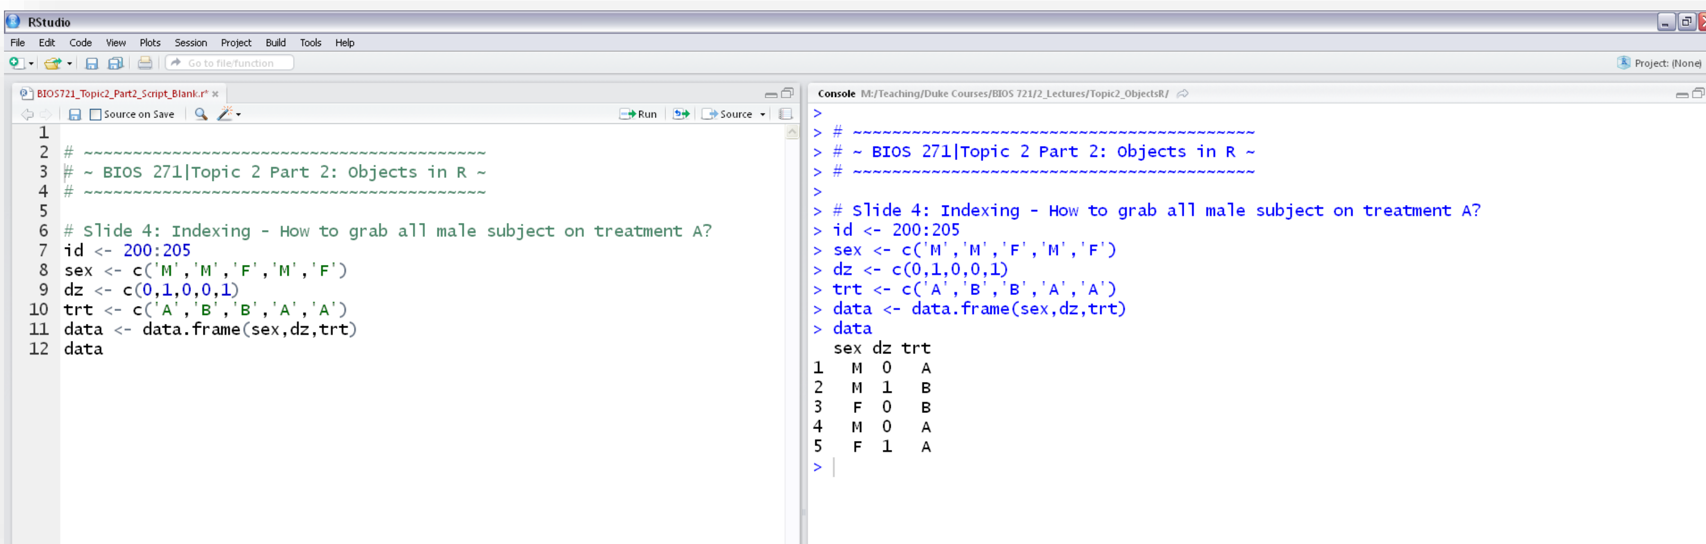

Indexing – Motivating Example

- Recall this example?
  - The code on the previous slide uses R operators to determine whether each observation satisfies a particular condition.
  - But this alone doesn't completely automate the indexing needed to create the data frame sub – need an R function that … does what?

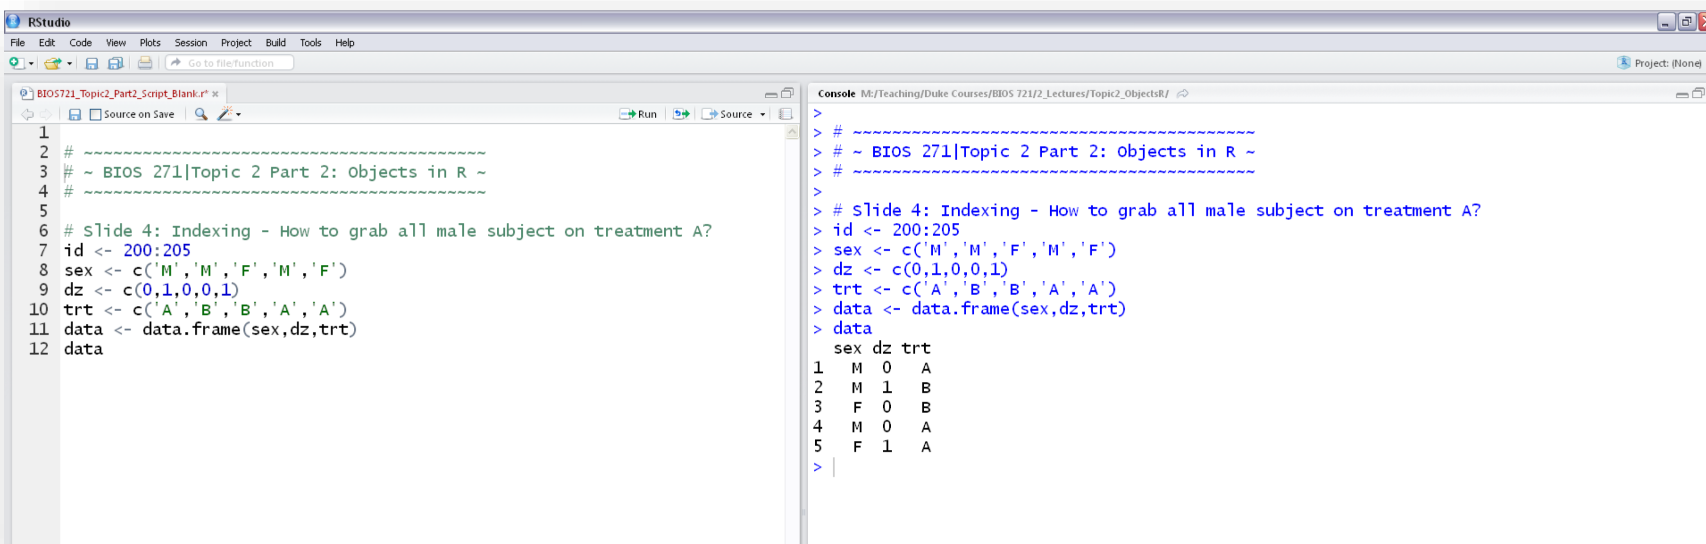

R Functions

- Where to begin?

- There a TON of R Functions.
  - Many are included in the base R installation
  - Many, many more that can be download in Packages developed by other R users
  - Only base R functions are “vetted”

- Today we will focus on simple R functions …
  - Ones that every statistician should have in their toolbox

R Functions

- The basic syntax associated with an R Function:
  - `function_name(inputs, …)`
  - The () tells R that the keyword being referenced is a function and not a data object (remember the [] tells R that an object is being indexed)

- To learn about functions, bring up the R help page by typing `?function_name` into the R console.
  - Example| `?mean`

- All functions belong to a package (look for `function_name{package_name}` in the upper left hand side of the R help page).
  - For non-base installation functions, you can download the package by using the `install.packages()` function and entering the name of the package in quotations as input.
  - We will discuss this more later in the course.

Useful R Functions

- Functions that will come in handy …
  - `sqrt(x)` → returns square root of x
  - `log(x)` → returns log base_e of x
  - `exp(x)` → returns e raised to the power x
  - `mean(x)` → returns mean/average of x
  - `median(x)` → returns median of x
  - `sd(x)` → returns standard deviation of x
  - `min(x)` → returns minimum of x
  - `max(x)` → returns maximum of x
  - `length(x)` → returns the number of elements in x
  - `dim(x)` → returns the no. of rows and cols of a matrix x
  - `solve(x)` → returns inverse of a matrix x
  - `t(x)` → returns transpose of a matrix x

Useful R Functions

- Functions that will come in handy …
  - `sqrt(x)` → returns
  - `log(x)` → returns
  - `exp(x)` → returns
  - `mean(x)` → returns
  - `median(x)` → returns median of x
  - `sd(x)` → returns standard deviation of x
  - `min(x)` → returns minimum of x
  - `max(x)` → returns maximum of x
  - `length(x)` → returns the number of elements in x
  - `dim(x)` → returns the no. of rows and cols of a matrix x
  - `solve(x)` → returns inverse of a matrix x
  - `t(x)` → returns transpose of a matrix x

When encountering new R functions, you should always ask …

- What type of object does it apply to?
- What type of object does it return?
- And how are these two objects related?

Useful R Functions

- Let's consider the function sqrt()
  - Example| `v <- 1:5; sqrt(v)`

    `m <- matrix(c(1,4,9,25),2,2); sqrt(m)`

- Now let's answer the 3 'function' questions:
  - What type of object does it apply to?
  - What type of object will it return?
  - How are these two objects related to each other?

Useful R Functions

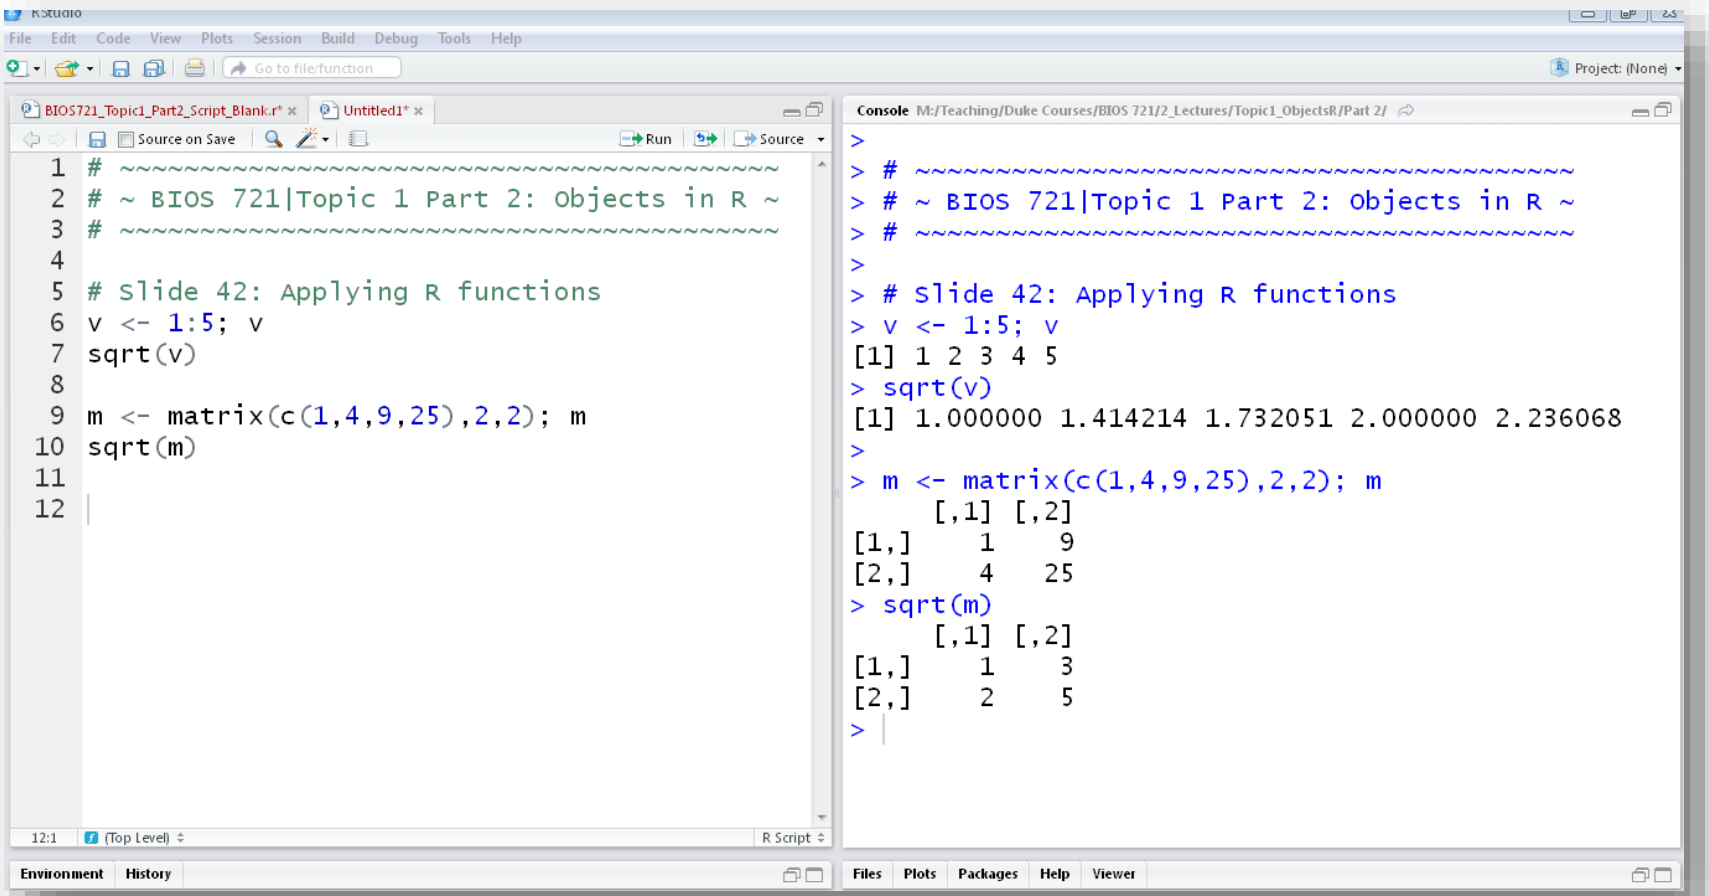

Useful R Functions

- Let's consider the function min()
  - Example| `v <- 1:5; min(v)`

    `m <- matrix(c(1,4,9,25),2,2); min(m)`

- Now let's answer the 3 'function' questions:
  - What type of object does it apply to?
  - What type of object will it return?
  - How are these two objects related to each other?

Useful R Functions

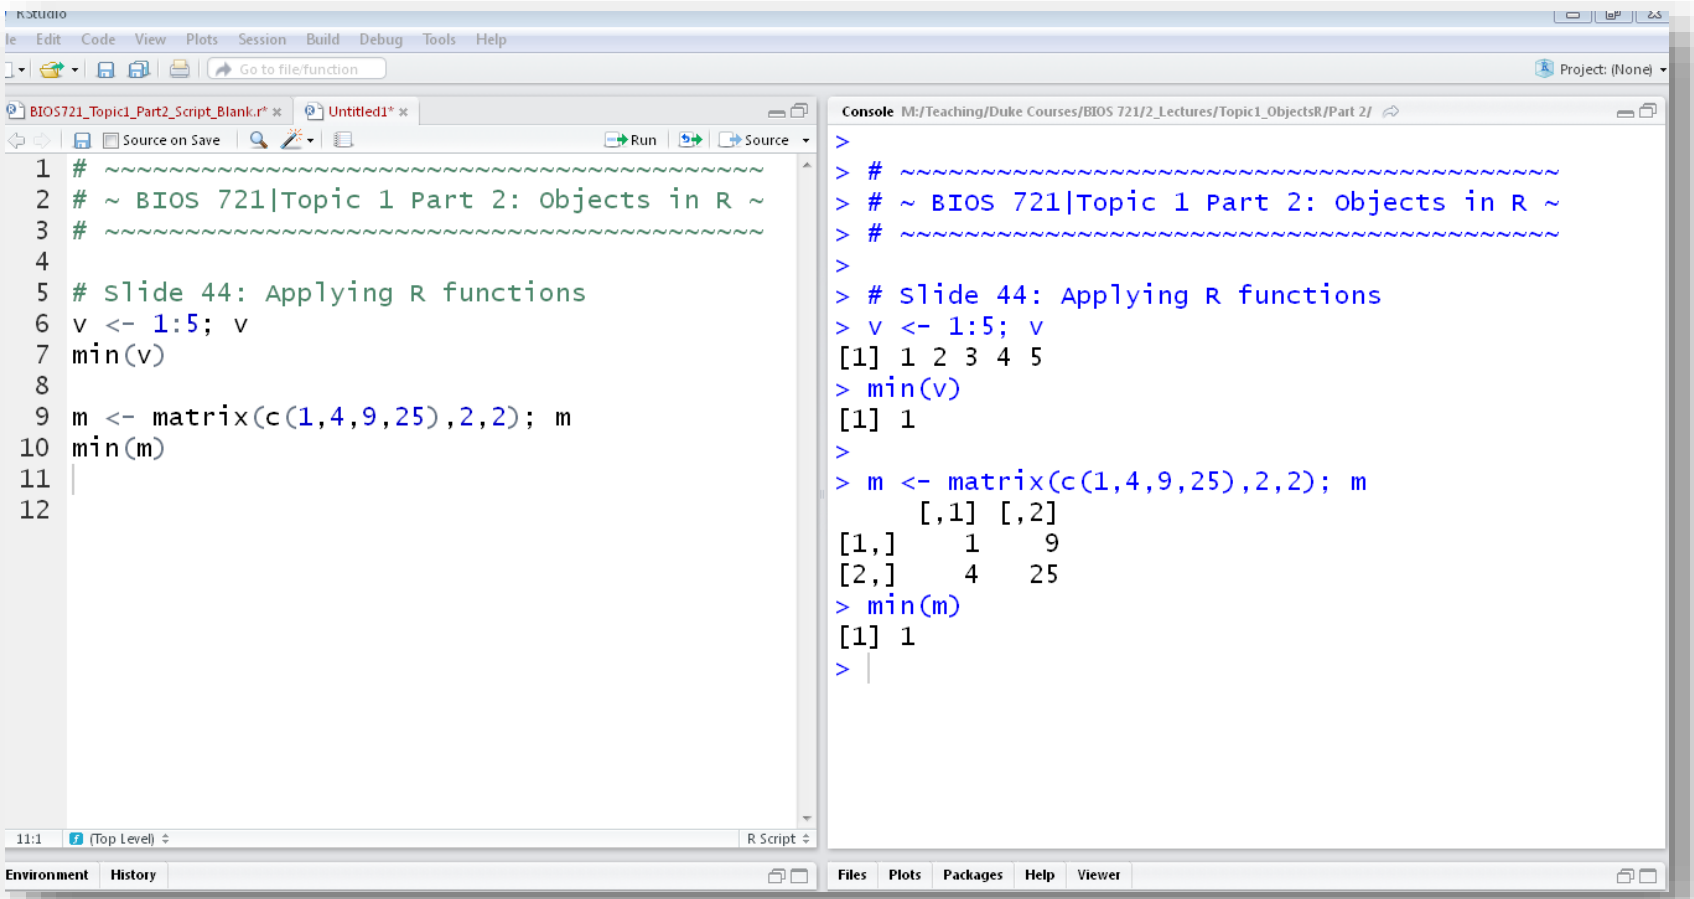

Useful R Functions

- Let's consider the function dim()
  - Example| `v <- 1:5; dim(v)`

    `m <- matrix(c(1,4,9,25),2,2); dim(m)`

- Now let's answer the 3 'function' questions:
  - What type of object does it apply to?
  - What type of object will it return?
  - How are these two objects related to each other?

Useful R Functions

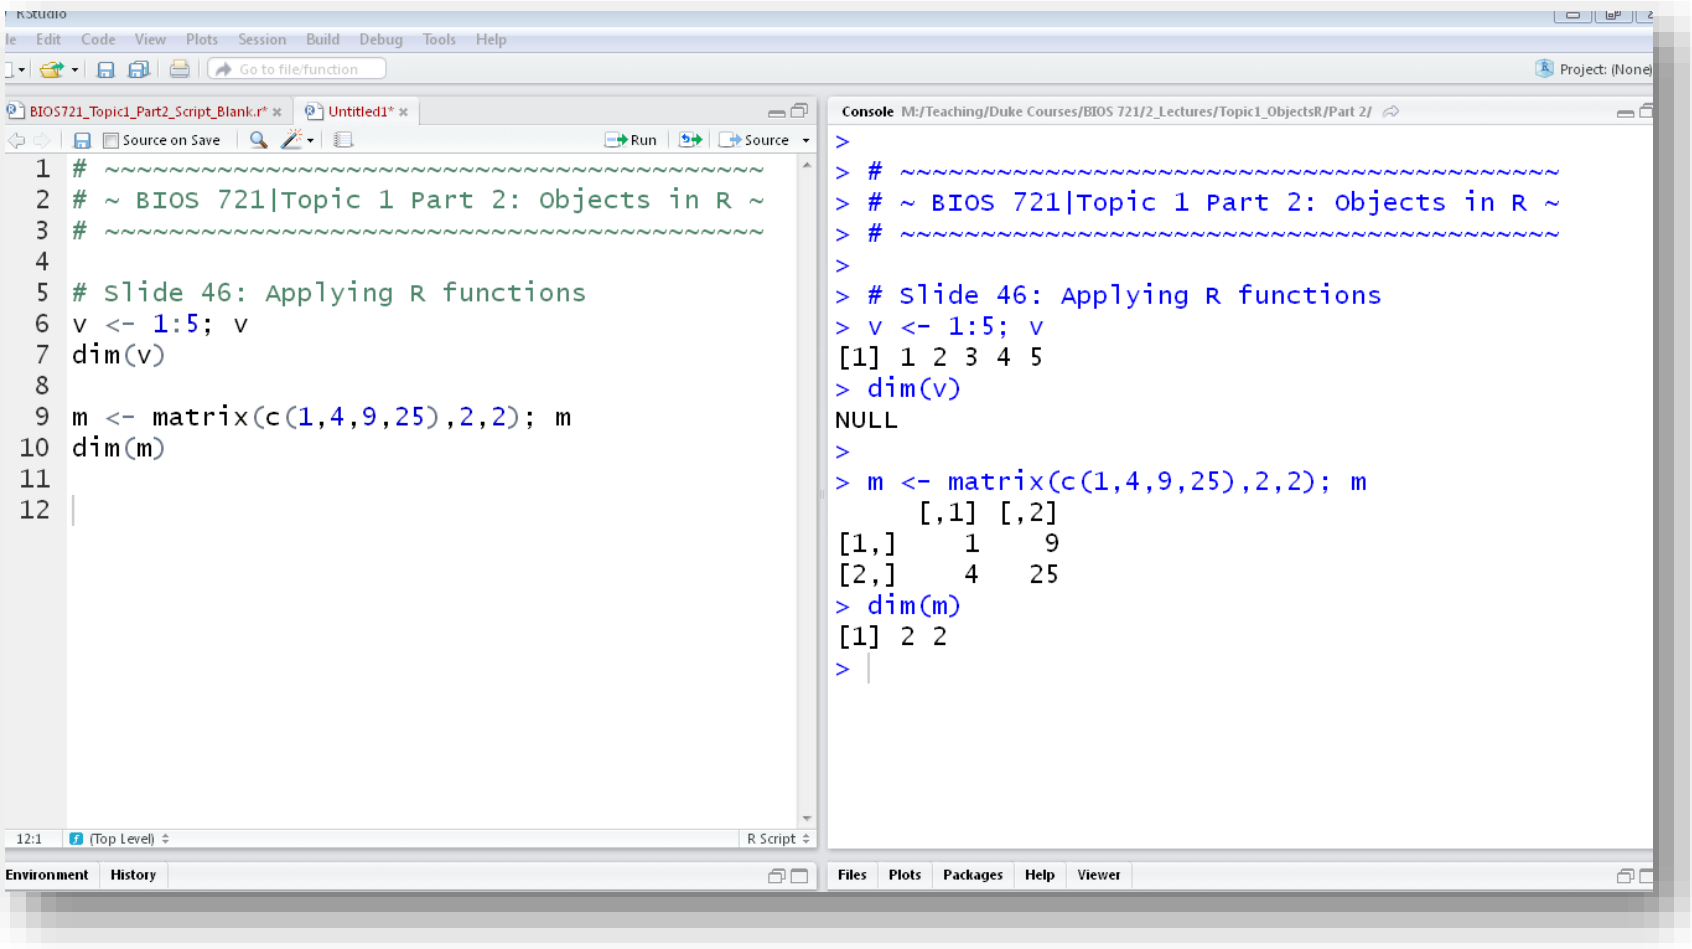

Useful R Functions

- Functions that will come in handy …
  - `sqrt(x)` → returns
  - `log(x)` → returns
  - `exp(x)` → returns
  - `mean(x)` → returns
  - `median(x)` → retu
  - `sd(x)` → returns sta
  - `min(x)` → returns
  - `max(x)` → returns maximum of x
  - `length(x)` → returns the number of elements in x
  - `dim(x)` → returns the no. of rows and cols of a matrix x
  - `solve(x)` → returns inverse of a matrix x
  - `t(x)` → returns transpose of a matrix x

When encountering new R functions, you should always ask …

- What type of object does it apply to?
- What type of object does it return?
- And how are these two objects related?

This is just a starter list … simple functions that all statisticians should have in their toolbox!

Other Useful R Functions

- To randomly shuffle/select elements in a vector use
  - `sample(x,size,replace=FALSE…)`
    - Returns `size` randomly selected elements from the vector `x`.
    - Can selected with (`replace=TRUE`) or without (`replace=FALSE`) replacement.

-
```
v <- -3:3
sample(v,4)
sample(v,7)
sample(v,7,replace=TRUE)
```

Other Useful R Functions

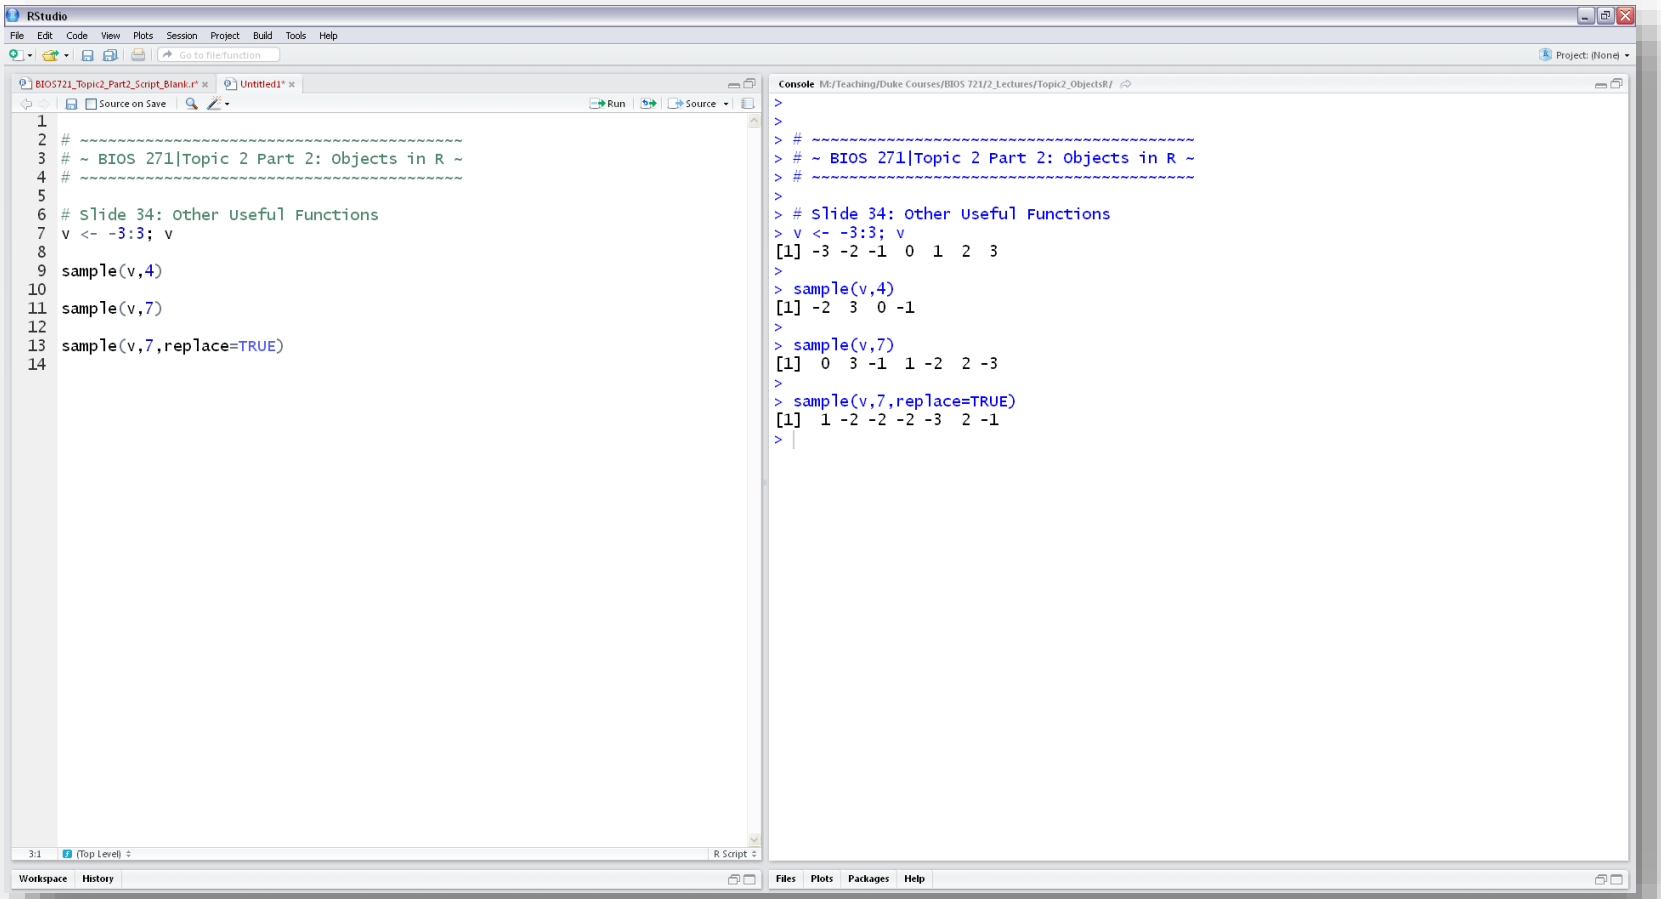

Other Useful R Functions

- Can use `which()` to determine which elements of a vector satisfy a condition
  - ```
    v <- c(1,5,4,0,9)
    which(v==4)
    which(v<=4)
    which(v>11)
    ```
- Can then use the results of `which()` to index a vector
  - ```
    id <- 20:24
    id[which(v==4)]
    id[which(v<=4)]
    id[which(v>11)]
    ```

Other Useful R Functions

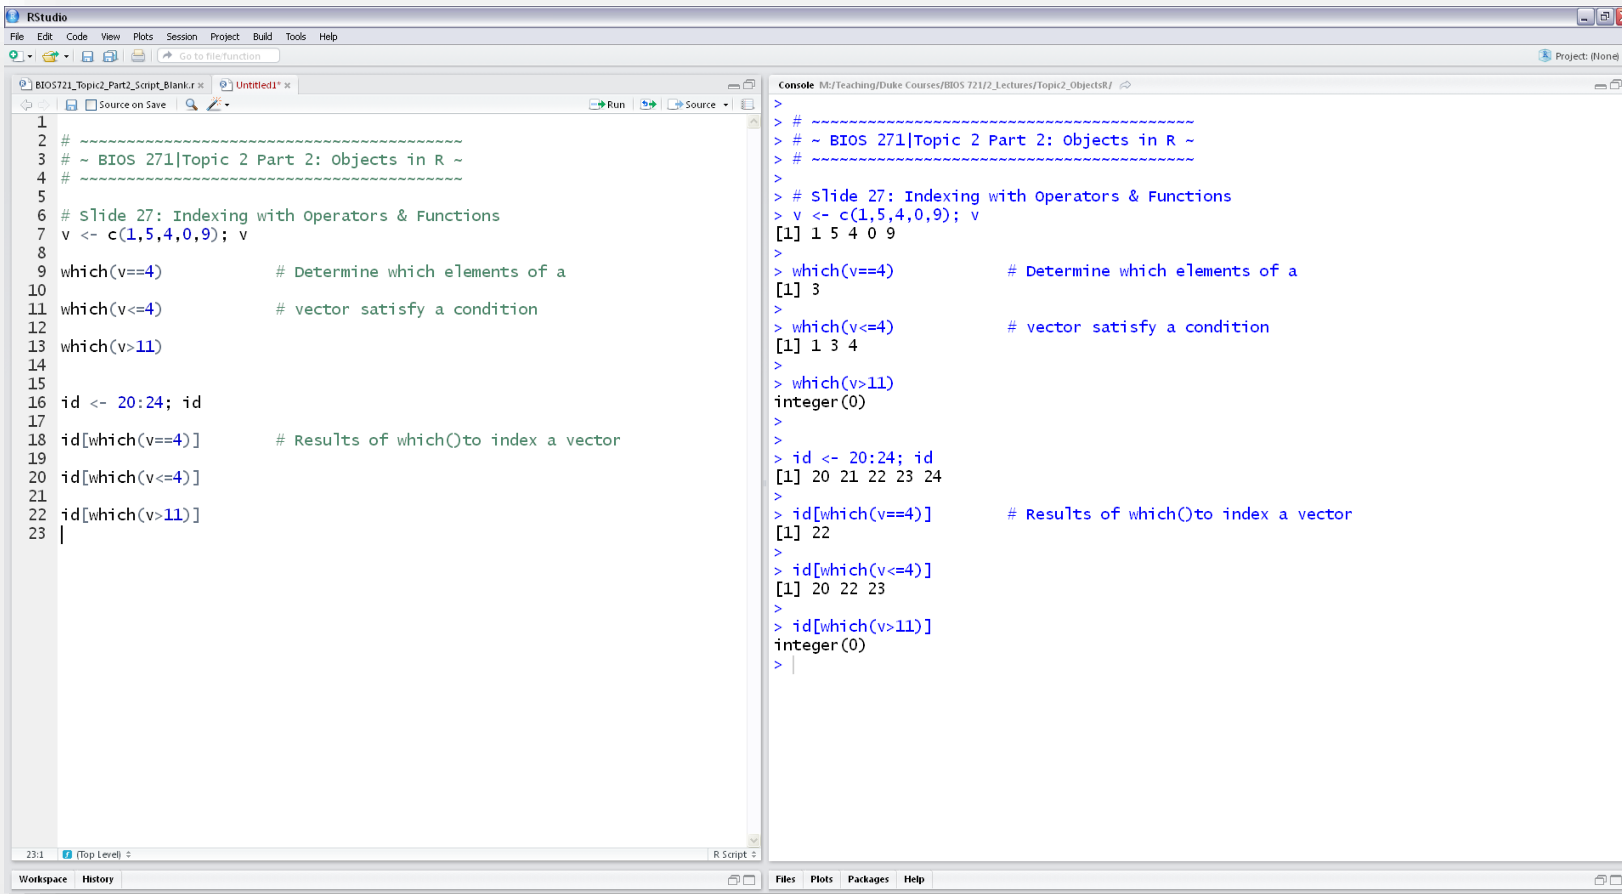

Putting it Together – Automated Indexing

- Back to this example …
  - Consider the data frame below. Suppose we wanted to create another data frame that subsets to all male subjects who were exposed to treatment A.
  - How can we achieve this using R in an automated fashion (i.e. for any size data frame)?

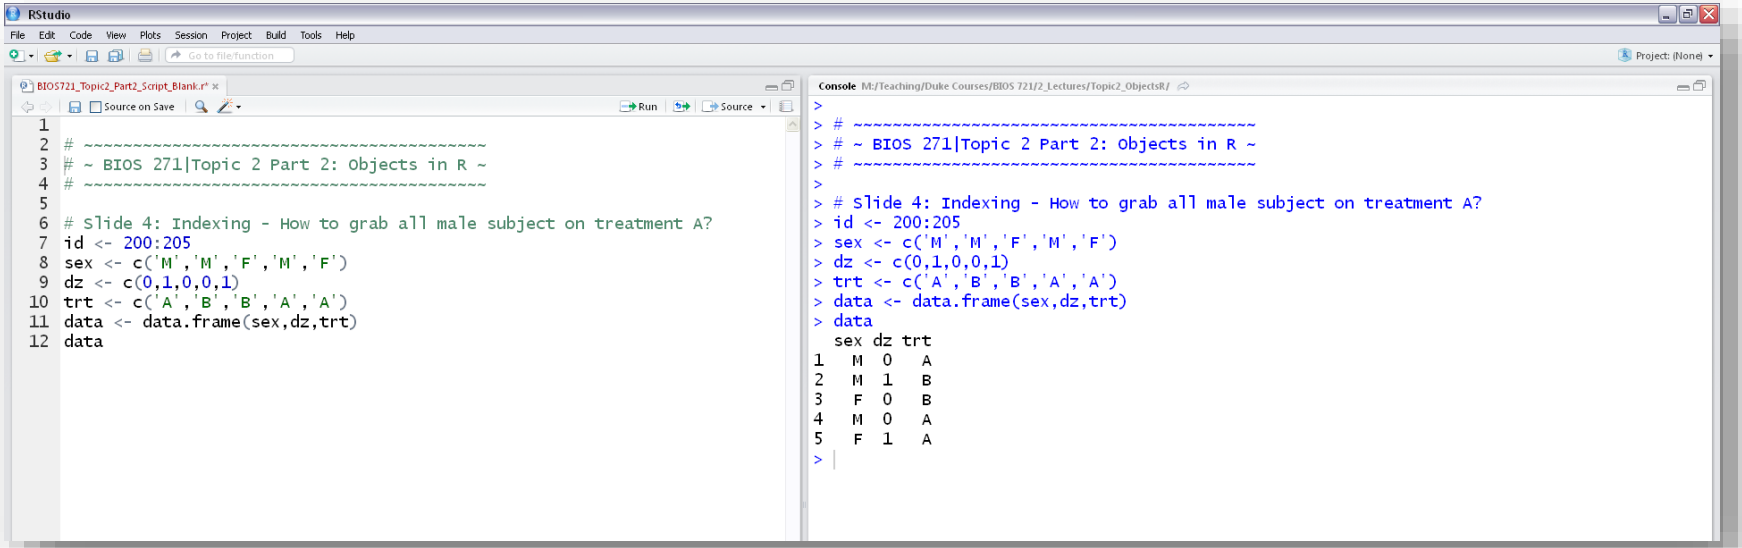

Indexing – Motivating Example

- Recall this example?
  - For solution – come to class!

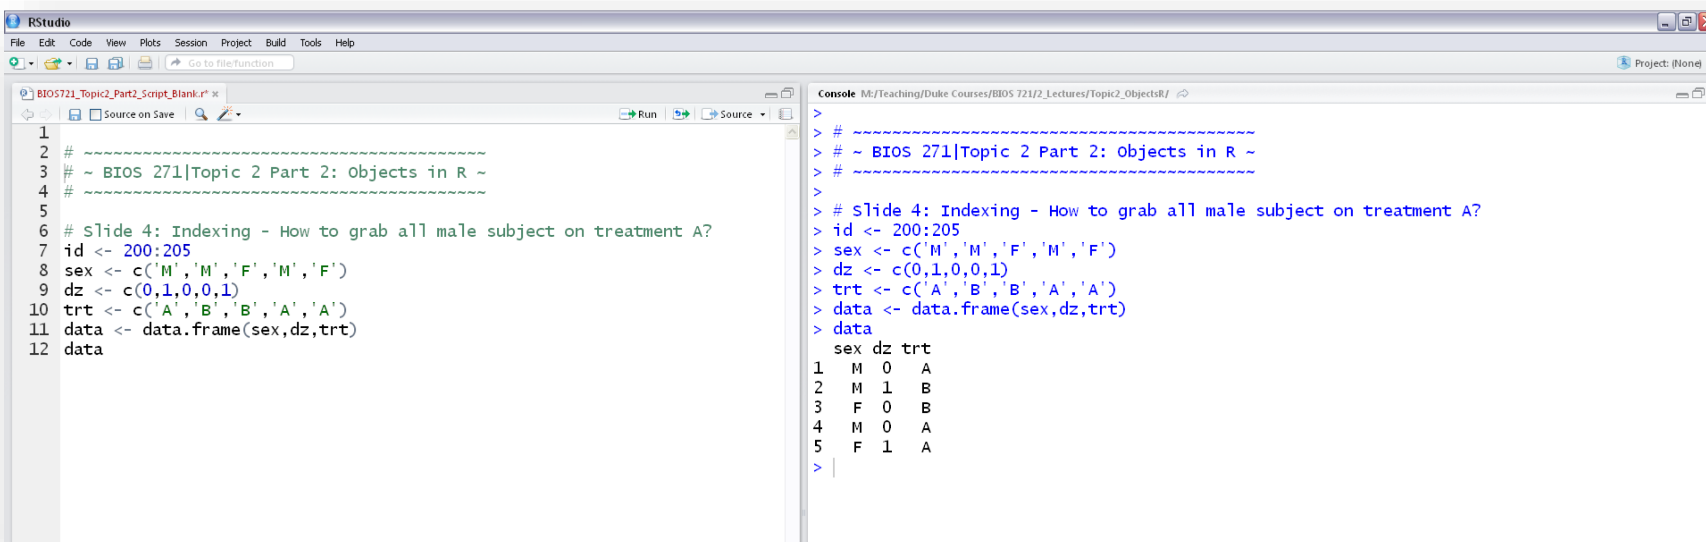

Next Time

- Next Time begin Topic 2
  - Data Management in R